

Show the **Right Numbers**

Data Visualization: Session 4

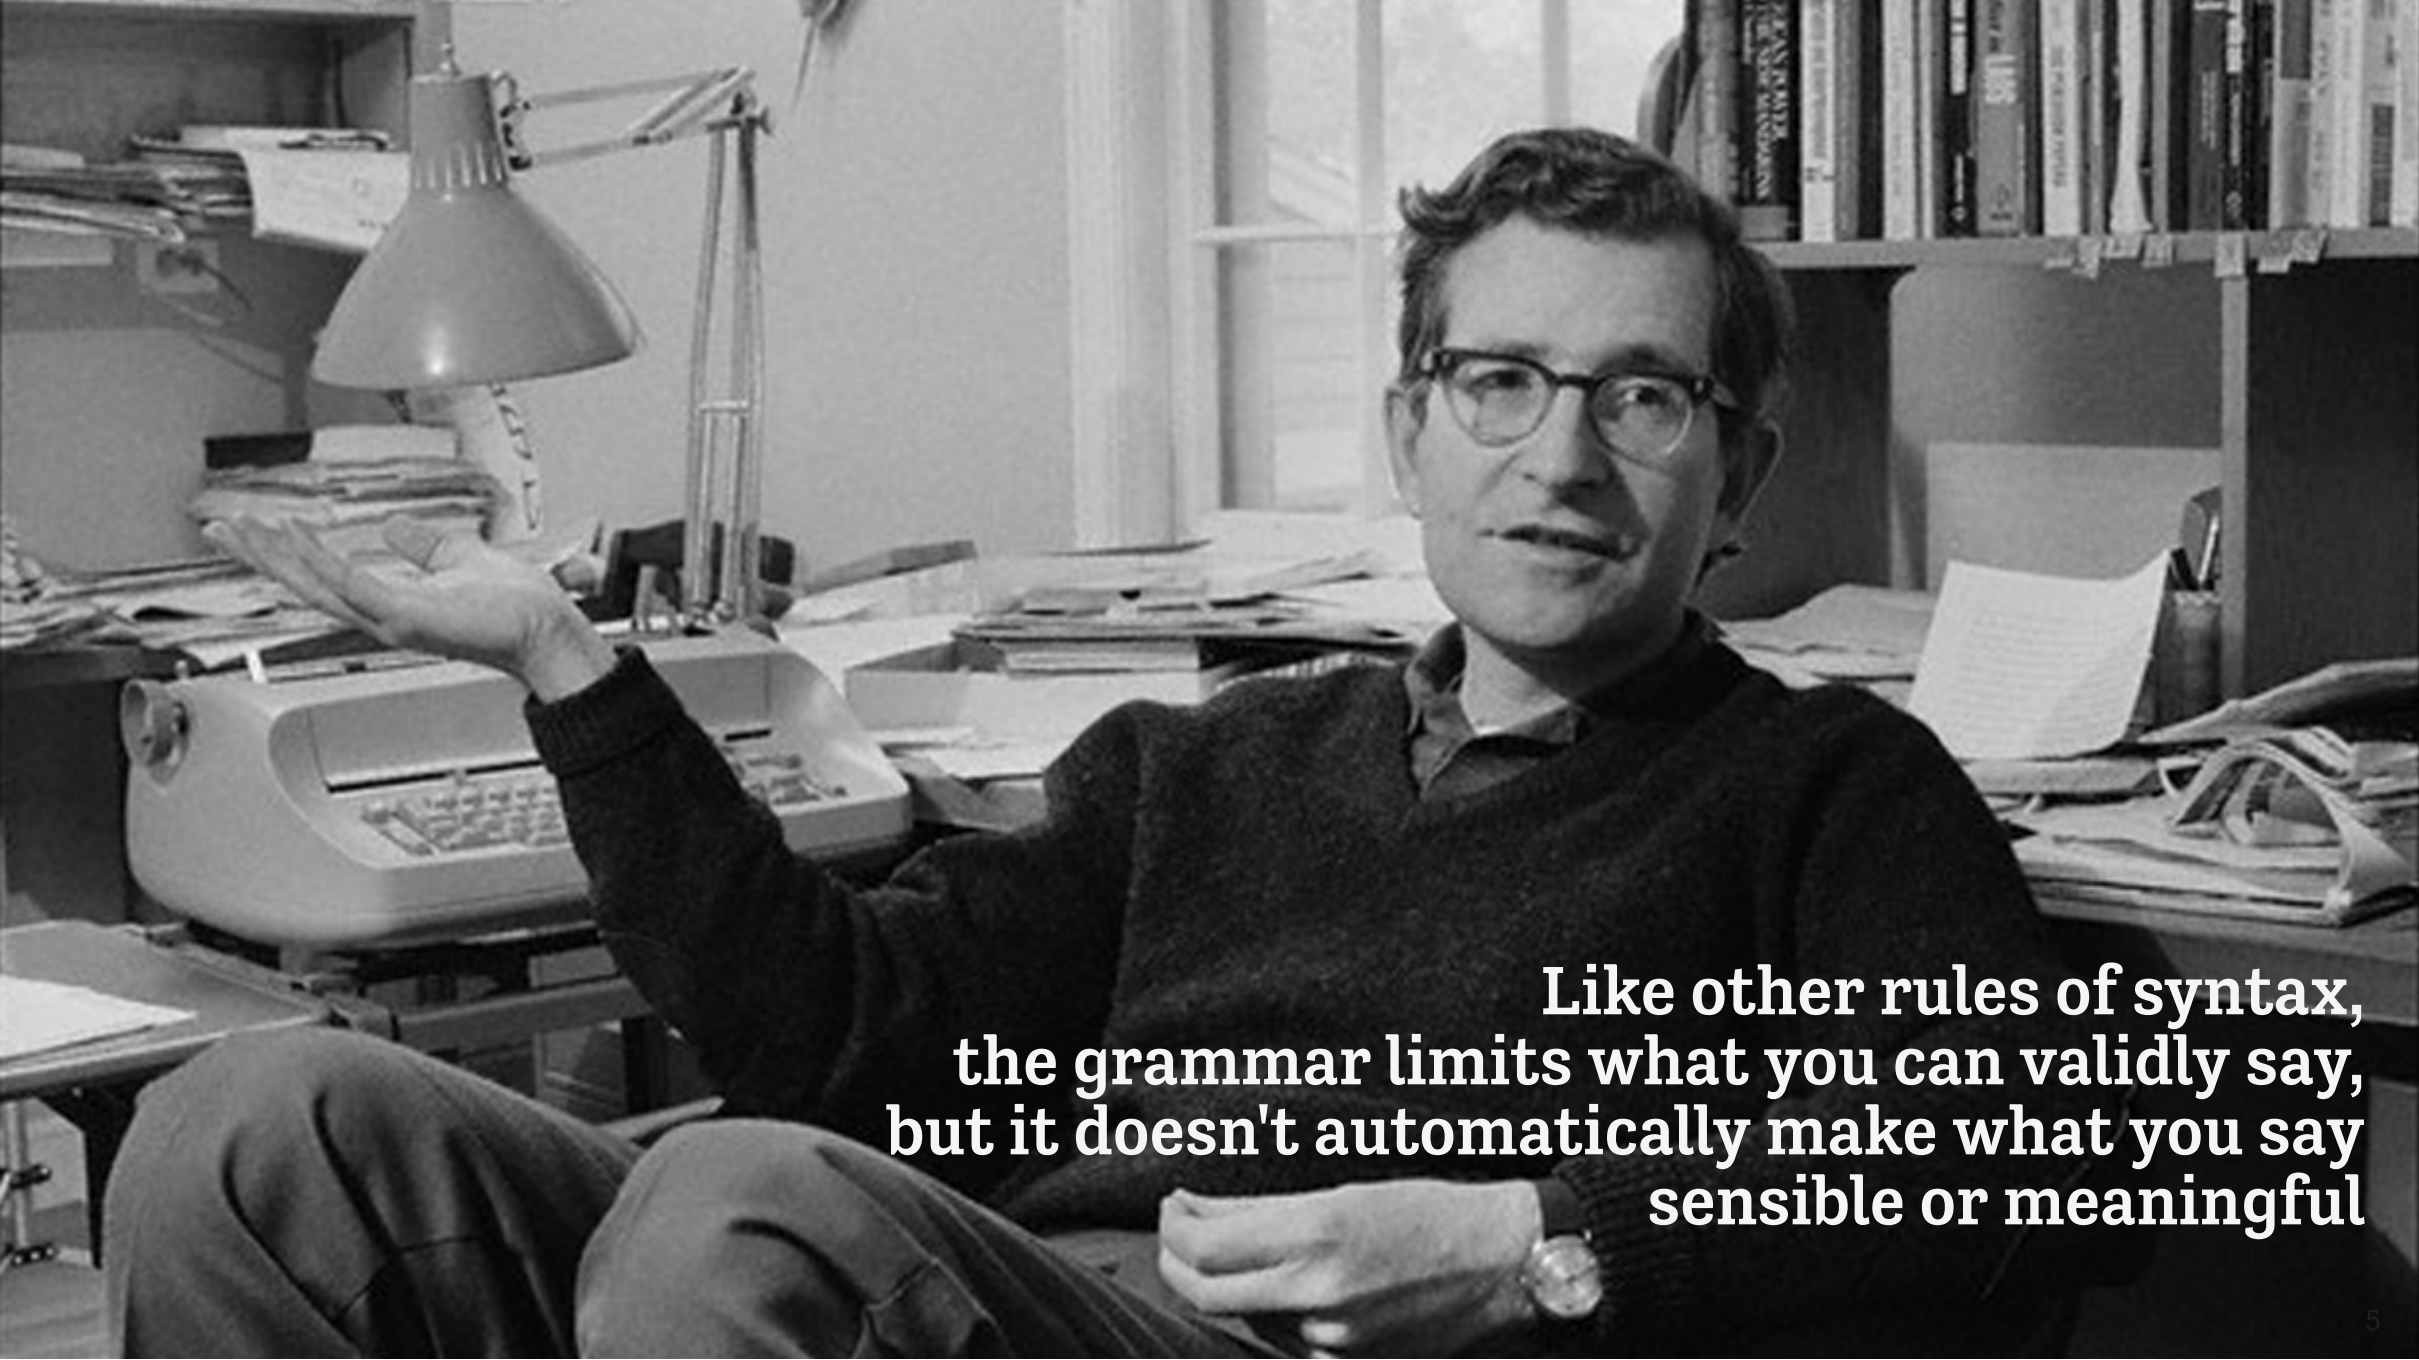
Kieran Healy
Code Horizons, April 2023

Set up our workspace

```
library(tidyverse)    # Your friend and mine  
library(gapminder)   # Gapminder data  
library(here)        # Portable file paths  
library(socviz)      # Handy socviz functions
```

ggplot
implements a
grammar of
graphics

The grammar is a set of rules for how to produce graphics from data, by *mapping* data to or *representing* it by geometric objects (like points and lines) that have aesthetic attributes (like position, color, size, and shape), together with further rules for transforming data if needed, for adjusting scales and their guides, and for projecting results onto some coordinate system.



**Like other rules of syntax,
the grammar limits what you can validly say,
but it doesn't automatically make what you say
sensible or meaningful**

Grouped data and the group aesthetic

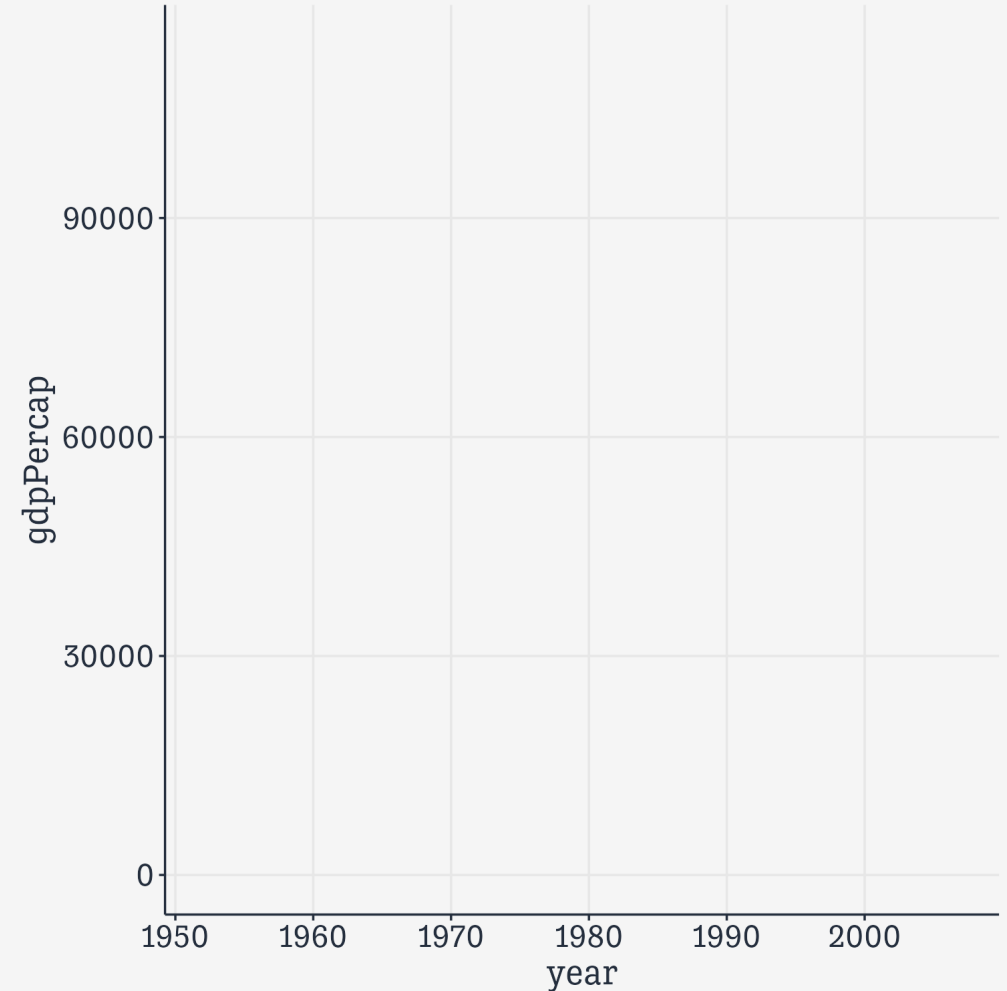
Try to make a lineplot

```
gapminder
```

```
## # A tibble: 1,704 × 6
##   country      continent year lifeExp      pop gdpPercap
##   <fct>        <fct>   <int>   <dbl>    <int>    <dbl>
## 1 Afghanistan Asia      1952    28.8  8425333    779.
## 2 Afghanistan Asia      1957    30.3  9240934    821.
## 3 Afghanistan Asia      1962    32.0 10267083    853.
## 4 Afghanistan Asia      1967    34.0 11537966    836.
## 5 Afghanistan Asia      1972    36.1 13079460    740.
## 6 Afghanistan Asia      1977    38.4 14880372    786.
## 7 Afghanistan Asia      1982    39.9 12881816    978.
## 8 Afghanistan Asia      1987    40.8 13867957    852.
## 9 Afghanistan Asia      1992    41.7 16317921    649.
## 10 Afghanistan Asia      1997    41.8 22227415    635.
## # i 1,694 more rows
```

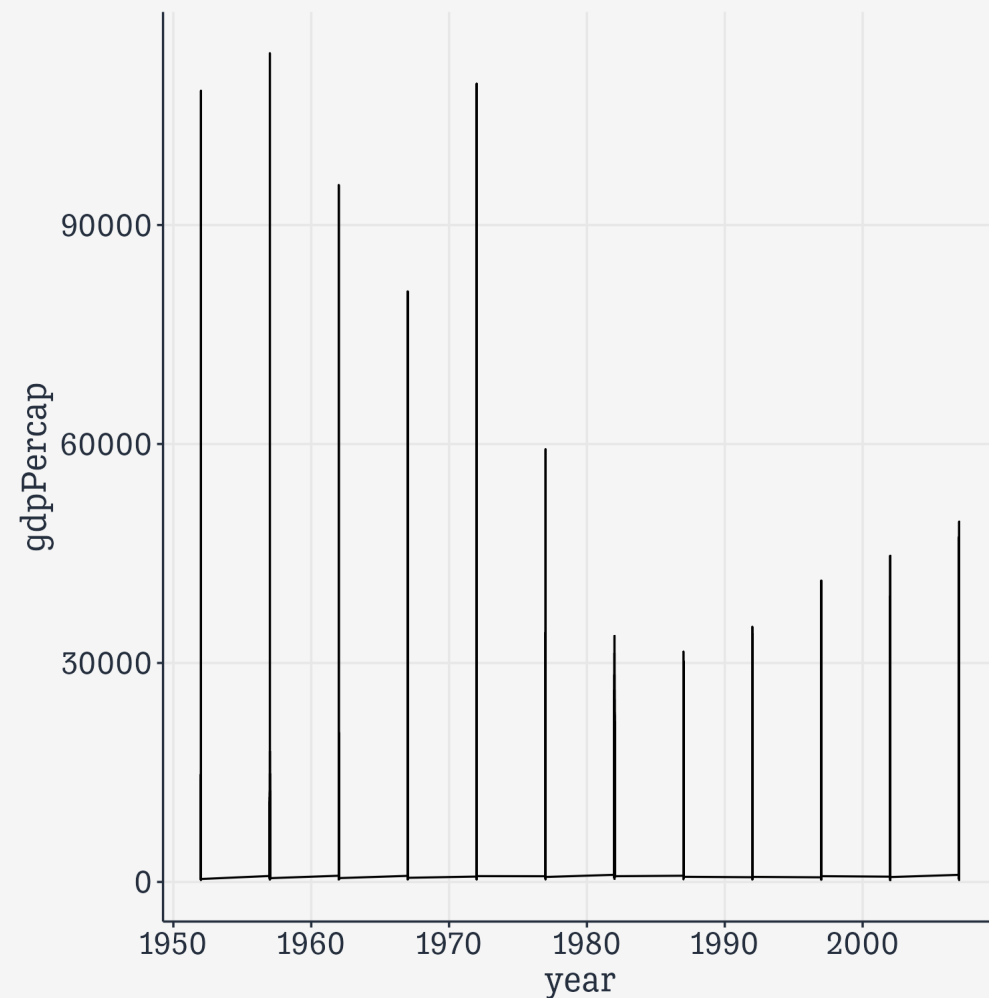
Try to make a lineplot

```
gapminder >  
ggplot(mapping = aes(x = year,  
                      y = gdpPercap))
```



Try to make a lineplot

```
gapminder >  
  ggplot(mapping = aes(x = year,  
                        y = gdpPercap)) +  
  geom_line()
```



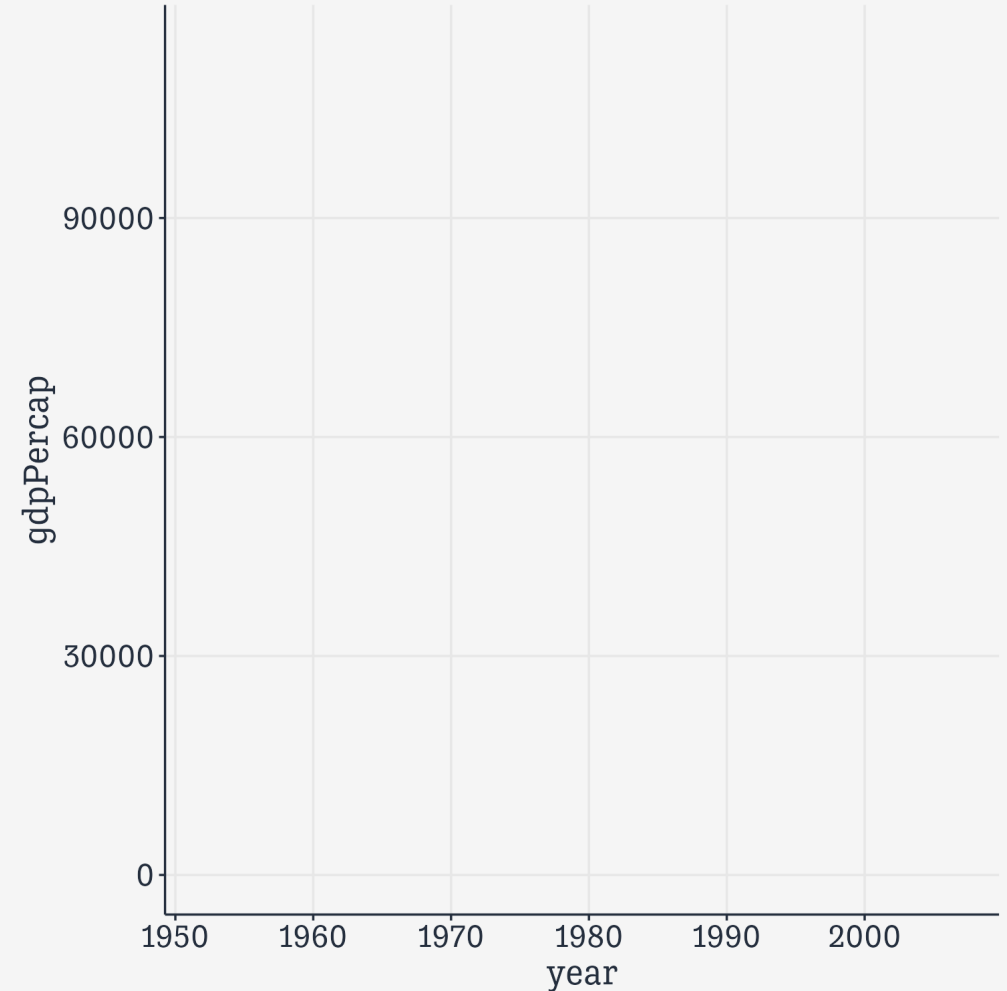
Try to make a lineplot

```
gapminder
```

```
## # A tibble: 1,704 × 6
##   country      continent  year lifeExp      pop gdpPercap
##   <fct>        <fct>    <int>  <dbl>    <int>    <dbl>
## 1 Afghanistan Asia      1952   28.8  8425333    779.
## 2 Afghanistan Asia      1957   30.3  9240934    821.
## 3 Afghanistan Asia      1962   32.0 10267083    853.
## 4 Afghanistan Asia      1967   34.0 11537966    836.
## 5 Afghanistan Asia      1972   36.1 13079460    740.
## 6 Afghanistan Asia      1977   38.4 14880372    786.
## 7 Afghanistan Asia      1982   39.9 12881816    978.
## 8 Afghanistan Asia      1987   40.8 13867957    852.
## 9 Afghanistan Asia      1992   41.7 16317921    649.
## 10 Afghanistan Asia      1997   41.8 22227415    635.
## # i 1,694 more rows
```

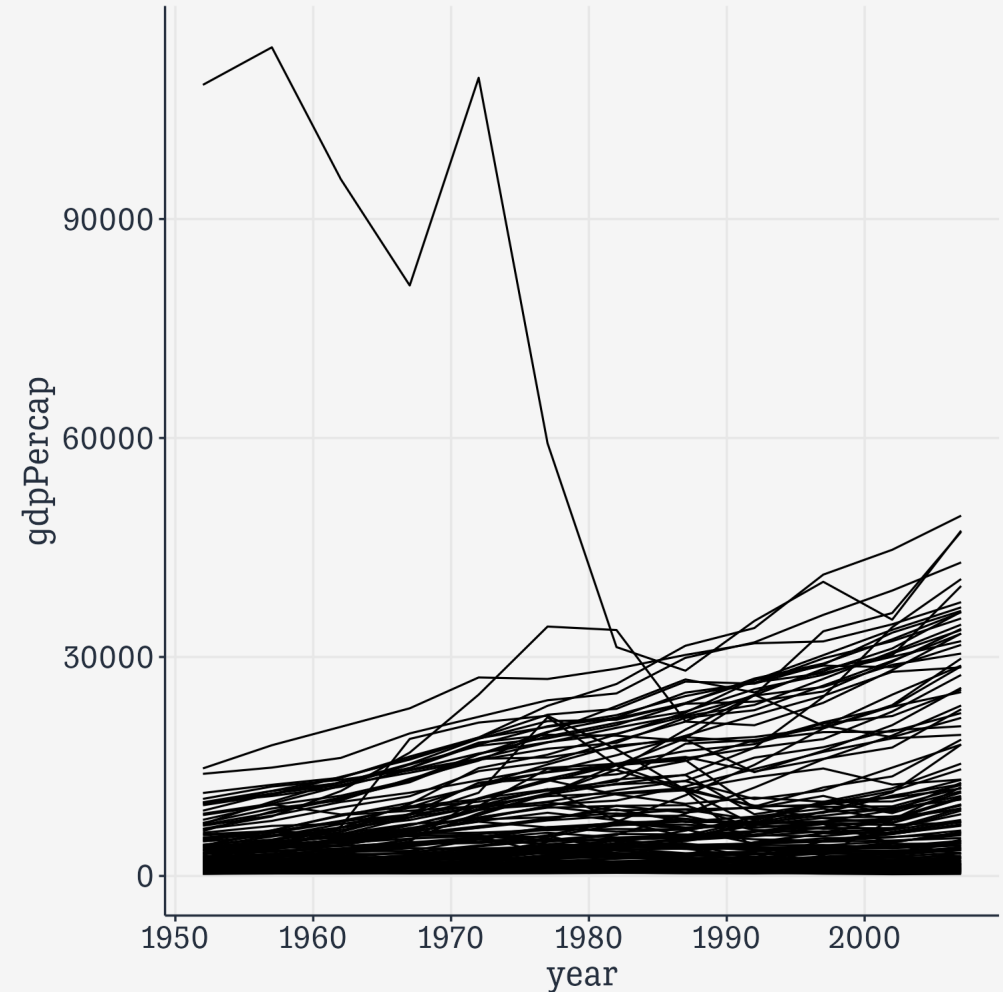
Try to make a lineplot

```
gapminder >  
ggplot(mapping = aes(x = year,  
                      y = gdpPercap))
```



Try to make a lineplot

```
gapminder >  
  ggplot(mapping = aes(x = year,  
                        y = gdpPercap)) +  
  geom_line(mapping = aes(group = country))
```



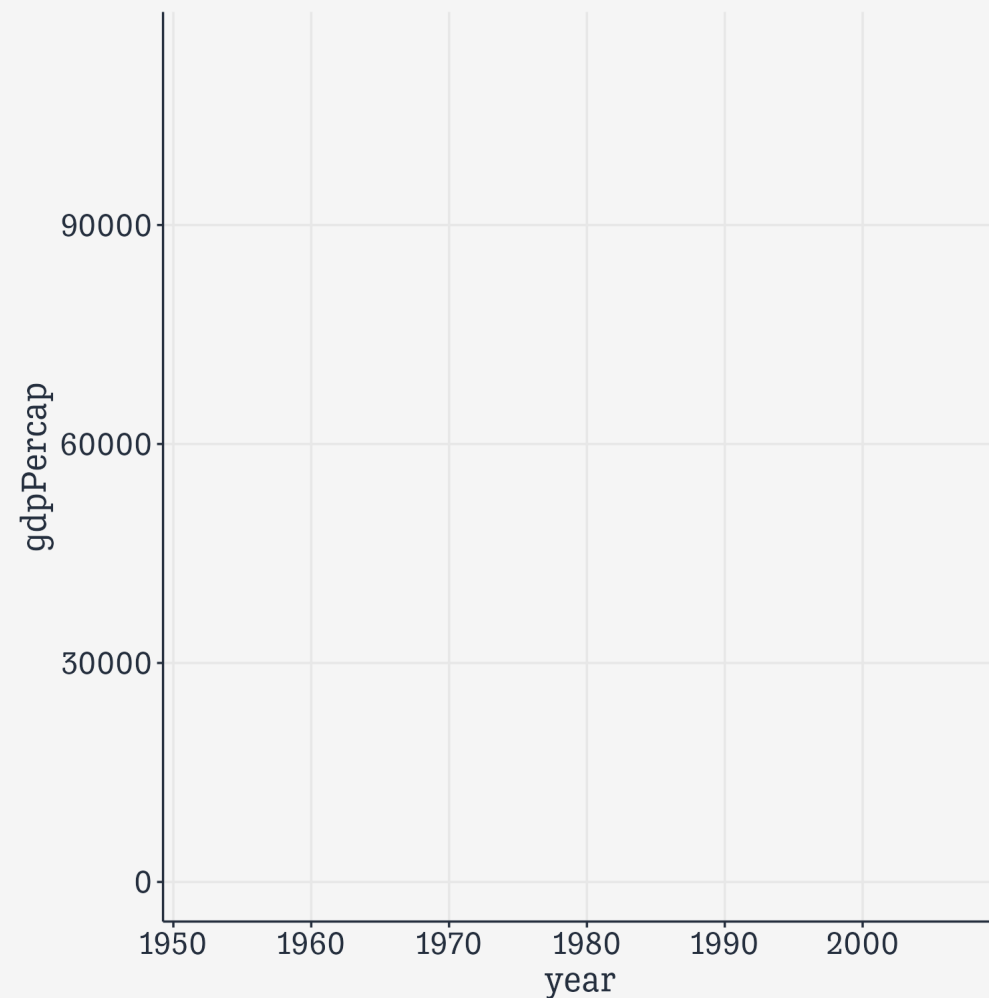
Facet the plot

gapminder

```
## # A tibble: 1,704 × 6
##   country      continent  year lifeExp      pop gdpPercap
##   <fct>        <fct>    <int>  <dbl>    <int>    <dbl>
## 1 Afghanistan Asia      1952   28.8  8425333    779.
## 2 Afghanistan Asia      1957   30.3  9240934    821.
## 3 Afghanistan Asia      1962   32.0 10267083    853.
## 4 Afghanistan Asia      1967   34.0 11537966    836.
## 5 Afghanistan Asia      1972   36.1 13079460    740.
## 6 Afghanistan Asia      1977   38.4 14880372    786.
## 7 Afghanistan Asia      1982   39.9 12881816    978.
## 8 Afghanistan Asia      1987   40.8 13867957    852.
## 9 Afghanistan Asia      1992   41.7 16317921    649.
## 10 Afghanistan Asia      1997   41.8 22227415    635.
## # i 1,694 more rows
```


Facet the plot

```
gapminder >  
  ggplot(mapping =  
    aes(x = year,  
        y = gdpPercap))
```



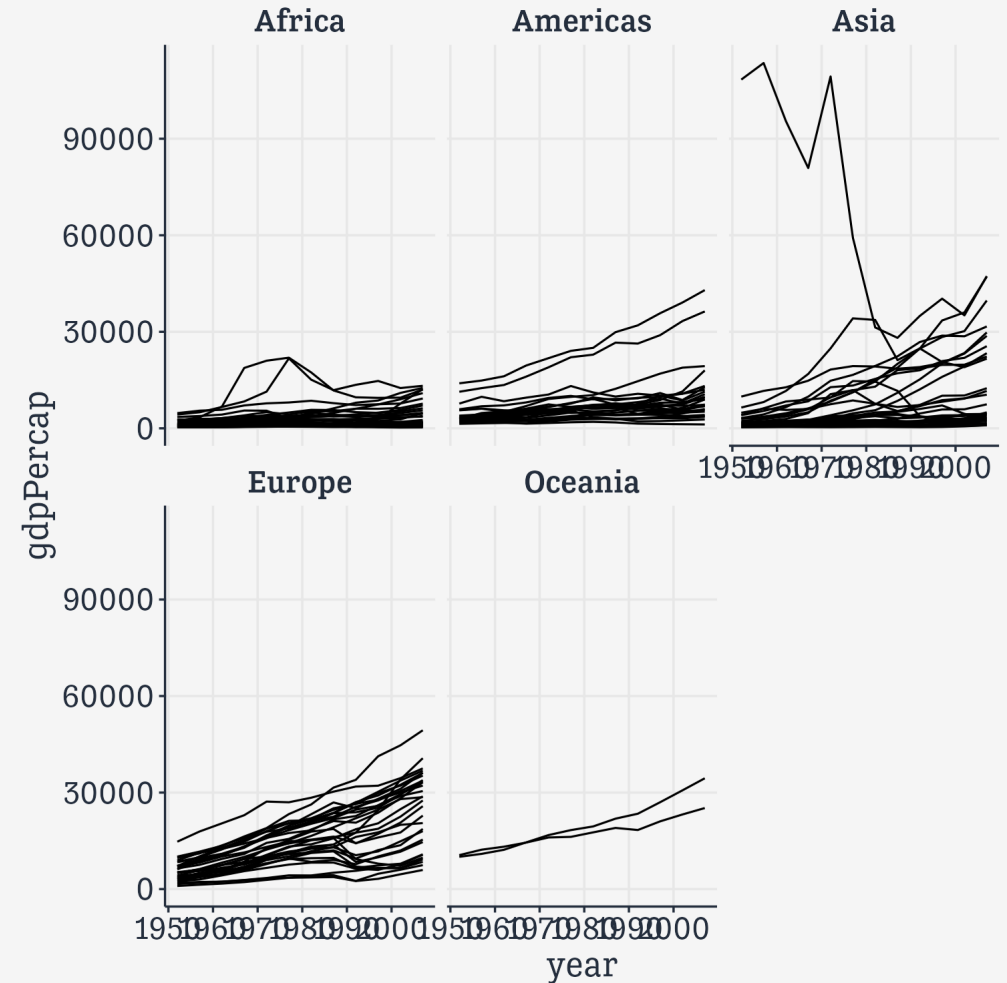
Facet the plot

```
gapminder >  
  ggplot(mapping =  
    aes(x = year,  
        y = gdpPercap)) +  
  geom_line(mapping = aes(group = country))
```



Facet the plot

```
gapminder >  
  ggplot(mapping =  
    aes(x = year,  
        y = gdpPercap)) +  
  geom_line(mapping = aes(group = country)) +  
  facet_wrap(~ continent)
```



Faceting is a very powerful tool

A facet is not a geom; it's a way of arranging repeated geoms by some additional variable

Faceting is a very powerful tool

A facet is not a geom; it's a way of arranging repeated geoms by some additional variable

Facets use R's "formula" syntax: `facet_wrap(~ continent)`

Faceting is a very powerful tool

A facet is not a geom; it's a way of arranging repeated geoms by some additional variable

Facets use R's "formula" syntax: `facet_wrap(~ continent)`

Read the `~` as "on" or "by"

Faceting is a very powerful tool

A facet is not a geom; it's a way of arranging repeated geoms by some additional variable

Facets use R's "formula" syntax: `facet_wrap(~ continent)`

Read the `~` as "on" or "by"

You can also use this syntax: `facet_wrap(vars(continent))`

Faceting is a very powerful tool

A facet is not a geom; it's a way of arranging repeated geoms by some additional variable

Facets use R's "formula" syntax: `facet_wrap(~ continent)`

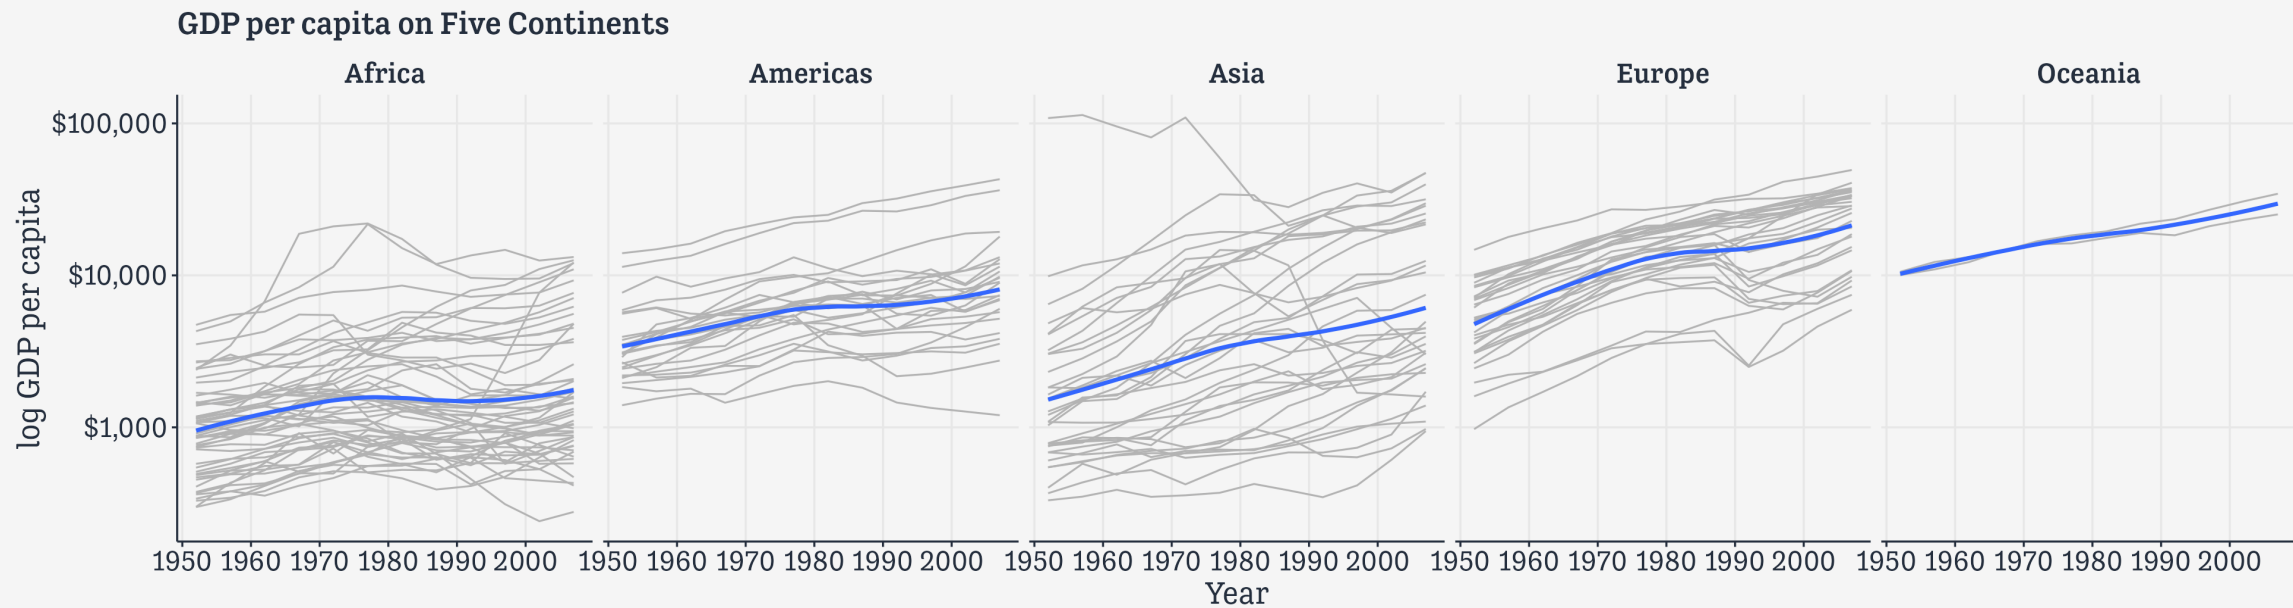
Read the `~` as "on" or "by"

You can also use this syntax: `facet_wrap(vars(continent))`

This is newer, and consistent with other ways of referring to variables within tidyverse functions.

Facets in action

```
p ← ggplot(data = gapminder,  
           mapping = aes(x = year,  
                         y = gdpPercap))  
  
p_out ← p + geom_line(color="gray70",  
                     mapping=aes(group = country)) +  
  geom_smooth(linewidth = 1.1,  
             method = "loess",  
             se = FALSE) +  
  scale_y_log10(labels=scales::label_dollar()) +  
  facet_wrap(~ continent, ncol = 5) +  
  labs(x = "Year",  
       y = "log GDP per capita",  
       title = "GDP per capita on Five Continents",  
       caption = "Data: Gapminder")
```



Data: Gapminder

A more polished faceted plot.

One-variable summaries

The **midwest** dataset

County-level census data for Midwestern U.S. Counties

```
midwest
```

```
## # A tibble: 437 × 28
##   PID county state area poptotal popdensity popwhite popblack popamerindian
##   <int> <chr> <chr> <dbl> <int> <dbl> <int> <int> <int>
## 1 561 ADAMS IL 0.052 66090 1271. 63917 1702 98
## 2 562 ALEXAN... IL 0.014 10626 759 7054 3496 19
## 3 563 BOND IL 0.022 14991 681. 14477 429 35
## 4 564 BOONE IL 0.017 30806 1812. 29344 127 46
## 5 565 BROWN IL 0.018 5836 324. 5264 547 14
## 6 566 BUREAU IL 0.05 35688 714. 35157 50 65
## 7 567 CALHOUN IL 0.017 5322 313. 5298 1 8
## 8 568 CARROLL IL 0.027 16805 622. 16519 111 30
## 9 569 CASS IL 0.024 13437 560. 13384 16 8
## 10 570 CHAMPA... IL 0.058 173025 2983. 146506 16559 331
## # i 427 more rows
## # i 19 more variables: popasian <int>, popother <int>, percwhite <dbl>,
## # percblack <dbl>, percamerindian <dbl>, percasian <dbl>, percother <dbl>,
## # popadults <int>, perchsd <dbl>, percollege <dbl>, percprof <dbl>,
## # poppovertyknown <int>, percpovertyknown <dbl>, percbelowpoverty <dbl>,
## # percchildbelowpovert <dbl>, percadultpoverty <dbl>,
## # percelderlypoverty <dbl>, inmetro <int>, category <chr>
```

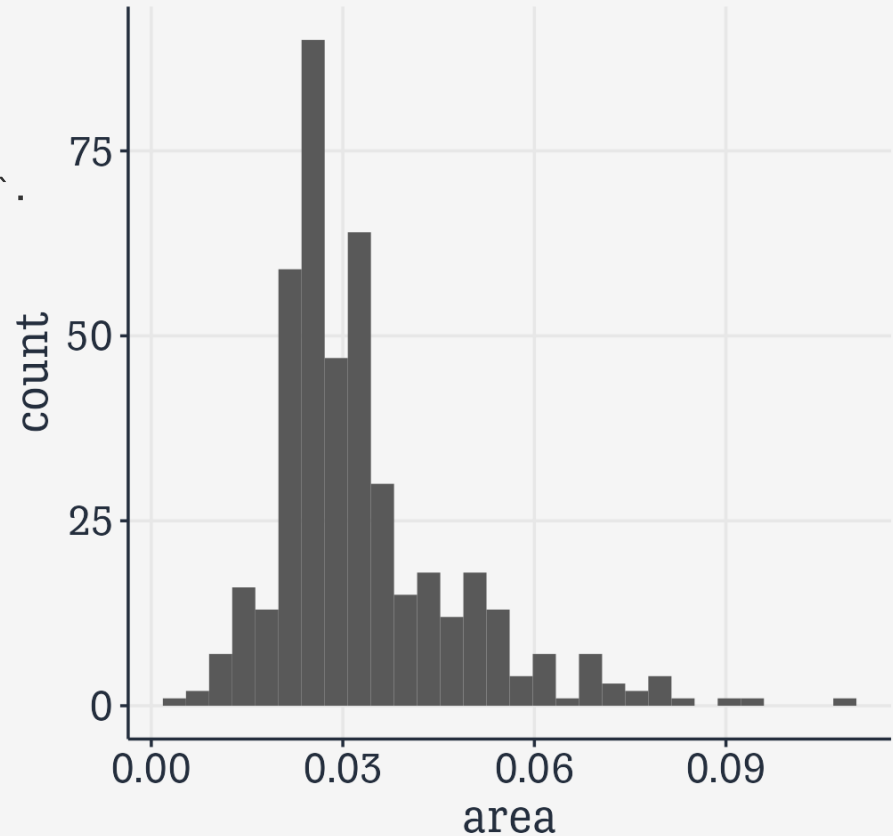
stat_ functions work behind the scenes

```
p <- ggplot(data = midwest,  
            mapping = aes(x = area))
```

```
p + geom_histogram()
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```


Here the default `stat_` function for this geom has to make a choice. It is letting us know we might want to override it.



stat_ functions work behind the scenes

```
p <- ggplot(data = midwest,  
            mapping = aes(x = area))  
  
p + geom_histogram(bins = 10)
```

We can choose *either* the number of bins *or* the binwidth



Compare two distributions

```
## Two state codes
oh_wi <- c("OH", "WI")

midwest >
  filter(state %in% oh_wi) >
  ggplot(mapping = aes(x = percollege,
                       fill = state)) +
  geom_histogram(alpha = 0.5,
                 position = "identity")
```

Here we do the whole thing in a **pipeline** using the pipe and the `dplyr` verb `filter()` to subset rows of the data by some condition.

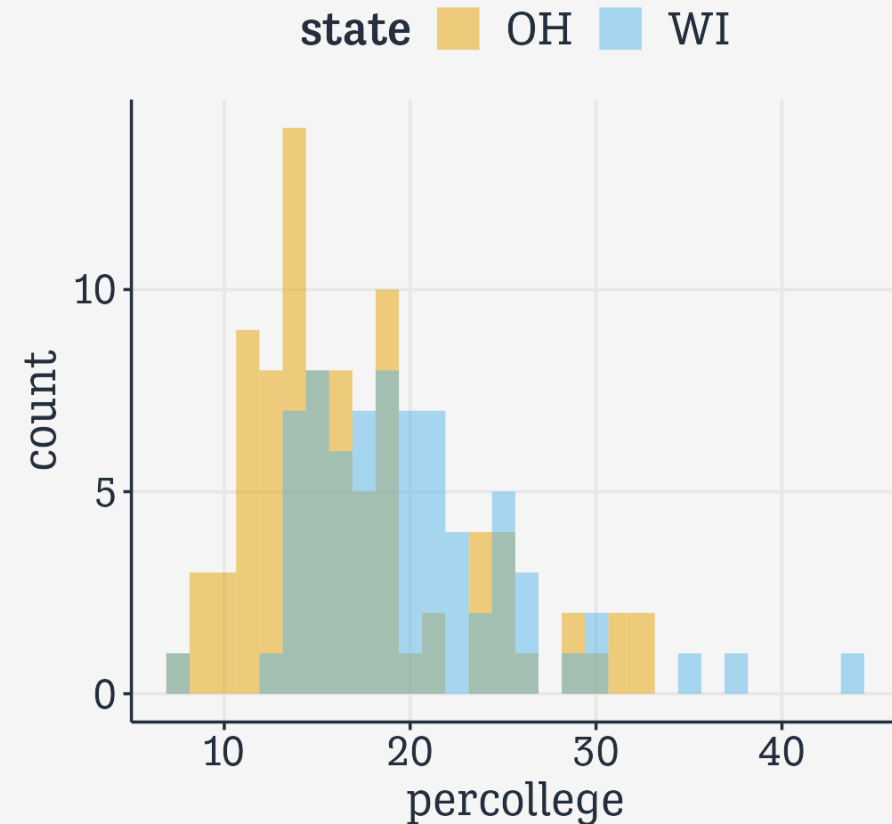
Experiment with leaving the `position` argument out, or changing it to `"dodge"`.

Compare two distributions

```
## Two state codes
oh_wi <- c("OH", "WI")

midwest >
  filter(state %in% oh_wi) >
  ggplot(mapping = aes(x = percollege,
                       fill = state)) +
  geom_histogram(alpha = 0.5,
                 position = "identity")
```

Here we do the whole thing in a **pipeline** using the pipe and the `dplyr` verb `filter()` to subset rows of the data by some condition. Experiment with leaving the `position` argument out, or changing it to "dodge".




geom_hist()'s counterpart, **geom_density()**

```
p ← ggplot(data = midwest,  
           mapping = aes(x = area))  
  
p + geom_density()
```

`geom_hist()`'s counterpart, `geom_density()`

```
p ← ggplot(data = midwest,  
            mapping = aes(x = area))  
  
p + geom_density()
```

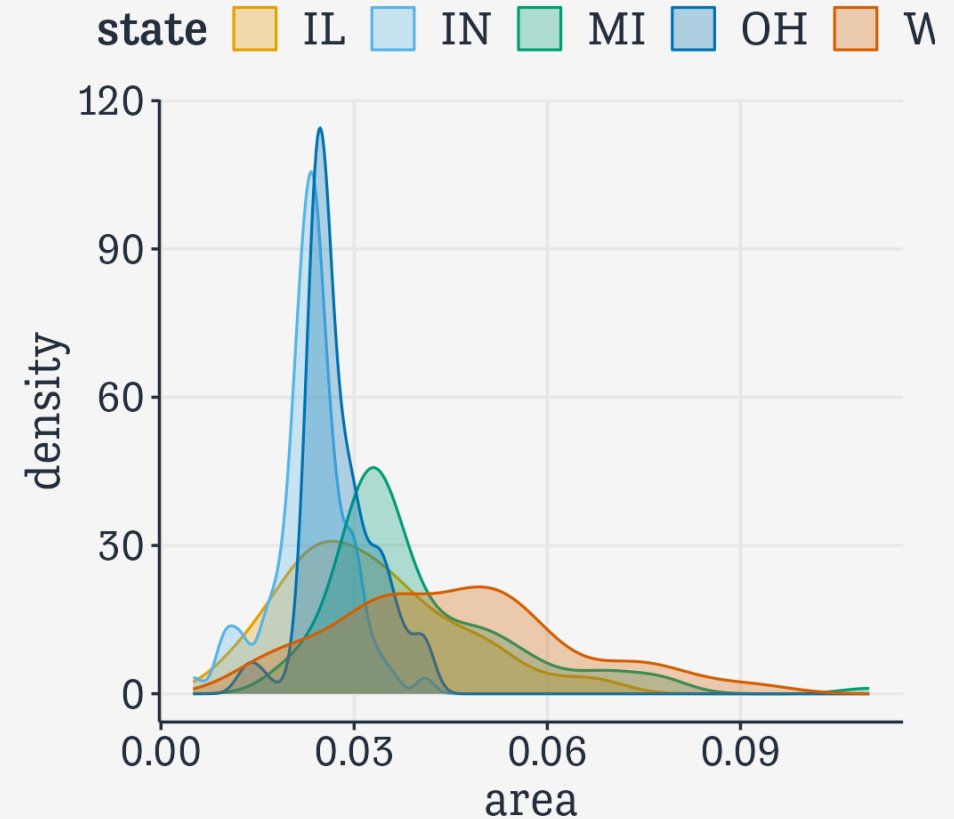


geom_hist()'s counterpart, geom_density()

```
p ← ggplot(data = midwest,  
  mapping = aes(x = area,  
    fill = state,  
    color = state))  
p + geom_density(alpha = 0.3)
```


geom_hist()'s counterpart, geom_density()

```
p <- ggplot(data = midwest,  
            mapping = aes(x = area,  
                          fill = state,  
                          color = state))  
p + geom_density(alpha = 0.3)
```



geom_hist()'s counterpart, geom_density()

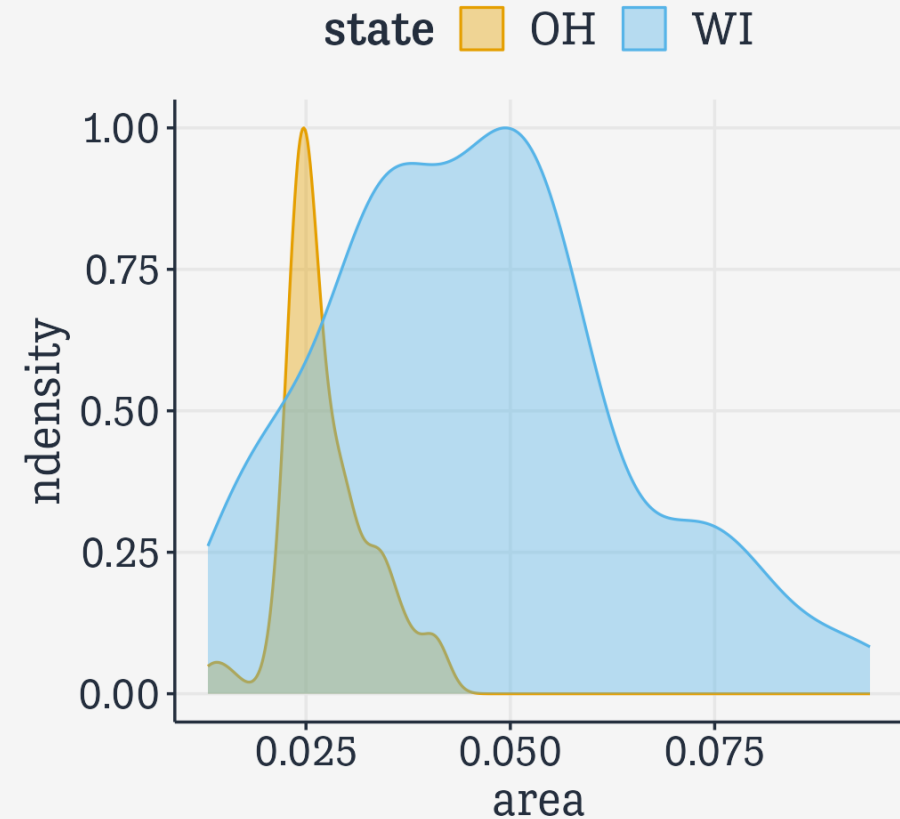
```
midwest >
  filter(state %in% oh_wi) >
  ggplot(mapping = aes(x = area,
                        fill = state,
                        color = state)) +
  geom_density(mapping = aes(y = after_stat(ndensity)),
              alpha = 0.4)
```

`after_stat(ndensity)` here is not in our data! It's *computed*. Histogram and density geoms have default statistics, but you can ask them to do more. The `stat_` functions associated with each `geom_` do this work behind the scenes.

geom_hist()'s counterpart, geom_density()

```
midwest >
  filter(state %in% oh_wi) >
  ggplot(mapping = aes(x = area,
                       fill = state,
                       color = state)) +
  geom_density(mapping = aes(y = after_stat(ndensity)),
              alpha = 0.4)
```

`after_stat(ndensity)` here is not in our data! It's *computed*. Histogram and density geoms have default statistics, but you can ask them to do more. The `stat_` functions associated with each `geom_` do this work behind the scenes.



**Compare subgroups to a
reference distribution**

Some made-up data

Consider 3,000 observations of some unit (e.g., a county) with summary measures for each group, and the population average.

```
df
## # A tibble: 3,000 × 5
##   unit  pop_a pop_b  pop_c pop_total
##   <int> <dbl> <dbl>   <dbl>   <dbl>
## 1     1  1.29  1.93 -0.0869    1.09
## 2     2  0.522 0.536 -0.762    0.190
## 3     3  2.14  1.47 -0.616    1.15
## 4     4  1.13  0.673 -0.242    0.575
## 5     5  1.04  1.30  1.18     1.12
## 6     6  1.80  0.140  2.05     1.33
## 7     7  0.186  1.30 -0.709    0.476
## 8     8 -0.953  0.520 -2.44    -0.767
## 9     9  0.700  1.66 -1.09     0.749
## 10    10  0.0416 0.484 -0.180    0.177
## # i 2,990 more rows
```

Get the data into long format!

df

```
## # A tibble: 3,000 × 5
##   unit  pop_a pop_b  pop_c pop_total
##   <int> <dbl> <dbl>  <dbl>    <dbl>
## 1     1  1.29  1.93 -0.0869    1.09
## 2     2  0.522  0.536 -0.762    0.190
## 3     3  2.14  1.47 -0.616    1.15
## 4     4  1.13  0.673 -0.242    0.575
## 5     5  1.04  1.30  1.18     1.12
## 6     6  1.80  0.140  2.05     1.33
## 7     7  0.186  1.30 -0.709    0.476
## 8     8 -0.953  0.520 -2.44    -0.767
## 9     9  0.700  1.66 -1.09     0.749
## 10    10  0.0416 0.484 -0.180    0.177
## # i 2,990 more rows
```

Get the data into long format!

```
df >
  pivot_longer(cols = pop_a:pop_total)

## # A tibble: 12,000 × 3
##   unit name      value
##   <int> <chr>    <dbl>
## 1     1 1 pop_a      1.29
## 2     1 1 pop_b      1.93
## 3     1 1 pop_c     -0.0869
## 4     1 1 pop_total  1.09
## 5     2 2 pop_a      0.522
## 6     2 2 pop_b      0.536
## 7     2 2 pop_c     -0.762
## 8     2 2 pop_total  0.190
## 9     3 3 pop_a      2.14
## 10    3 3 pop_b      1.47
## # i 11,990 more rows
```

First effort: Hard to read


```
df >
  pivot_longer(cols = pop_a:pop_total) >
  ggplot() +
  geom_histogram(mapping = aes(x = value,
                              y = after_stat(ncount),
                              color = name, fill = name),
                 stat = "bin", bins = 20,
                 size = 0.5, alpha = 0.7,
                 position = "identity") +
  labs(x = "Measure", y = "Scaled Count", color = "Group",
       fill = "Group",
       title = "Comparing Subgroups: Histograms")
```

Again, `after_stat(ncount)` is computed. The periods on either side are just a naming convention to show that the measure is computed by the `stat_` function (and to make sure it doesn't clash with any actual names in your data.)

First effort: Hard to read

```
df >
  pivot_longer(cols = pop_a:pop_total) >
  ggplot() +
  geom_histogram(mapping = aes(x = value,
                              y = after_stat(ncount),
                              color = name, fill = name),
                 stat = "bin", bins = 20,
                 size = 0.5, alpha = 0.7,
                 position = "identity") +
  labs(x = "Measure", y = "Scaled Count", color = "Group",
       fill = "Group",
       title = "Comparing Subgroups: Histograms")
```

Again, `after_stat(ncount)` is computed. The periods on either side are just a naming convention to show that the measure is computed by the `stat_` function (and to make sure it doesn't clash with any actual names in your data.)



A little pivot trick

```
# Treat pop_a to pop_total as a single variable  
df
```

```
## # A tibble: 3,000 × 5  
##   unit  pop_a pop_b  pop_c pop_total  
##   <int> <dbl> <dbl>   <dbl>   <dbl>  
## 1     1  1.29  1.93 -0.0869    1.09  
## 2     2  0.522 0.536 -0.762    0.190  
## 3     3  2.14  1.47 -0.616    1.15  
## 4     4  1.13  0.673 -0.242    0.575  
## 5     5  1.04  1.30  1.18     1.12  
## 6     6  1.80  0.140  2.05     1.33  
## 7     7  0.186  1.30 -0.709    0.476  
## 8     8 -0.953  0.520 -2.44    -0.767  
## 9     9  0.700  1.66 -1.09     0.749  
## 10    10  0.0416 0.484 -0.180    0.177  
## # i 2,990 more rows
```

A little pivot trick

```
# Treat pop_a to pop_total as a single variable  
df ►  
  pivot_longer(cols = pop_a:pop_total)
```

```
## # A tibble: 12,000 × 3  
##   unit name      value  
##   <int> <chr>    <dbl>  
## 1     1 pop_a      1.29  
## 2     1 pop_b      1.93  
## 3     1 pop_c     -0.0869  
## 4     1 pop_total  1.09  
## 5     2 pop_a      0.522  
## 6     2 pop_b      0.536  
## 7     2 pop_c     -0.762  
## 8     2 pop_total  0.190  
## 9     3 pop_a      2.14  
## 10    3 pop_b      1.47  
## # i 11,990 more rows
```

A little pivot trick

```
# Just treat pop_a to pop_c as the single variable.  
# Notice that pop_total just gets repeated.  
df
```

```
## # A tibble: 3,000 × 5  
##   unit  pop_a pop_b  pop_c pop_total  
##   <int> <dbl> <dbl> <dbl> <dbl>  
## 1     1  1.29  1.93 -0.0869  1.09  
## 2     2  0.522  0.536 -0.762  0.190  
## 3     3  2.14  1.47 -0.616  1.15  
## 4     4  1.13  0.673 -0.242  0.575  
## 5     5  1.04  1.30  1.18  1.12  
## 6     6  1.80  0.140  2.05  1.33  
## 7     7  0.186  1.30 -0.709  0.476  
## 8     8 -0.953  0.520 -2.44 -0.767  
## 9     9  0.700  1.66 -1.09  0.749  
## 10    10  0.0416  0.484 -0.180  0.177  
## # i 2,990 more rows
```

A little pivot trick

```
# Just treat pop_a to pop_c as the single variable.  
# Notice that pop_total just gets repeated.
```

```
df ▸
```

```
  pivot_longer(cols = pop_a:pop_c)
```

```
## # A tibble: 9,000 × 4  
##   unit pop_total name    value  
##   <int>    <dbl> <chr>    <dbl>  
## 1     1      1.09 pop_a     1.29  
## 2     1      1.09 pop_b     1.93  
## 3     1      1.09 pop_c    -0.0869  
## 4     2      0.190 pop_a     0.522  
## 5     2      0.190 pop_b     0.536  
## 6     2      0.190 pop_c    -0.762  
## 7     3      1.15 pop_a     2.14  
## 8     3      1.15 pop_b     1.47  
## 9     3      1.15 pop_c    -0.616  
## 10    4      0.575 pop_a     1.13  
## # i 8,990 more rows
```

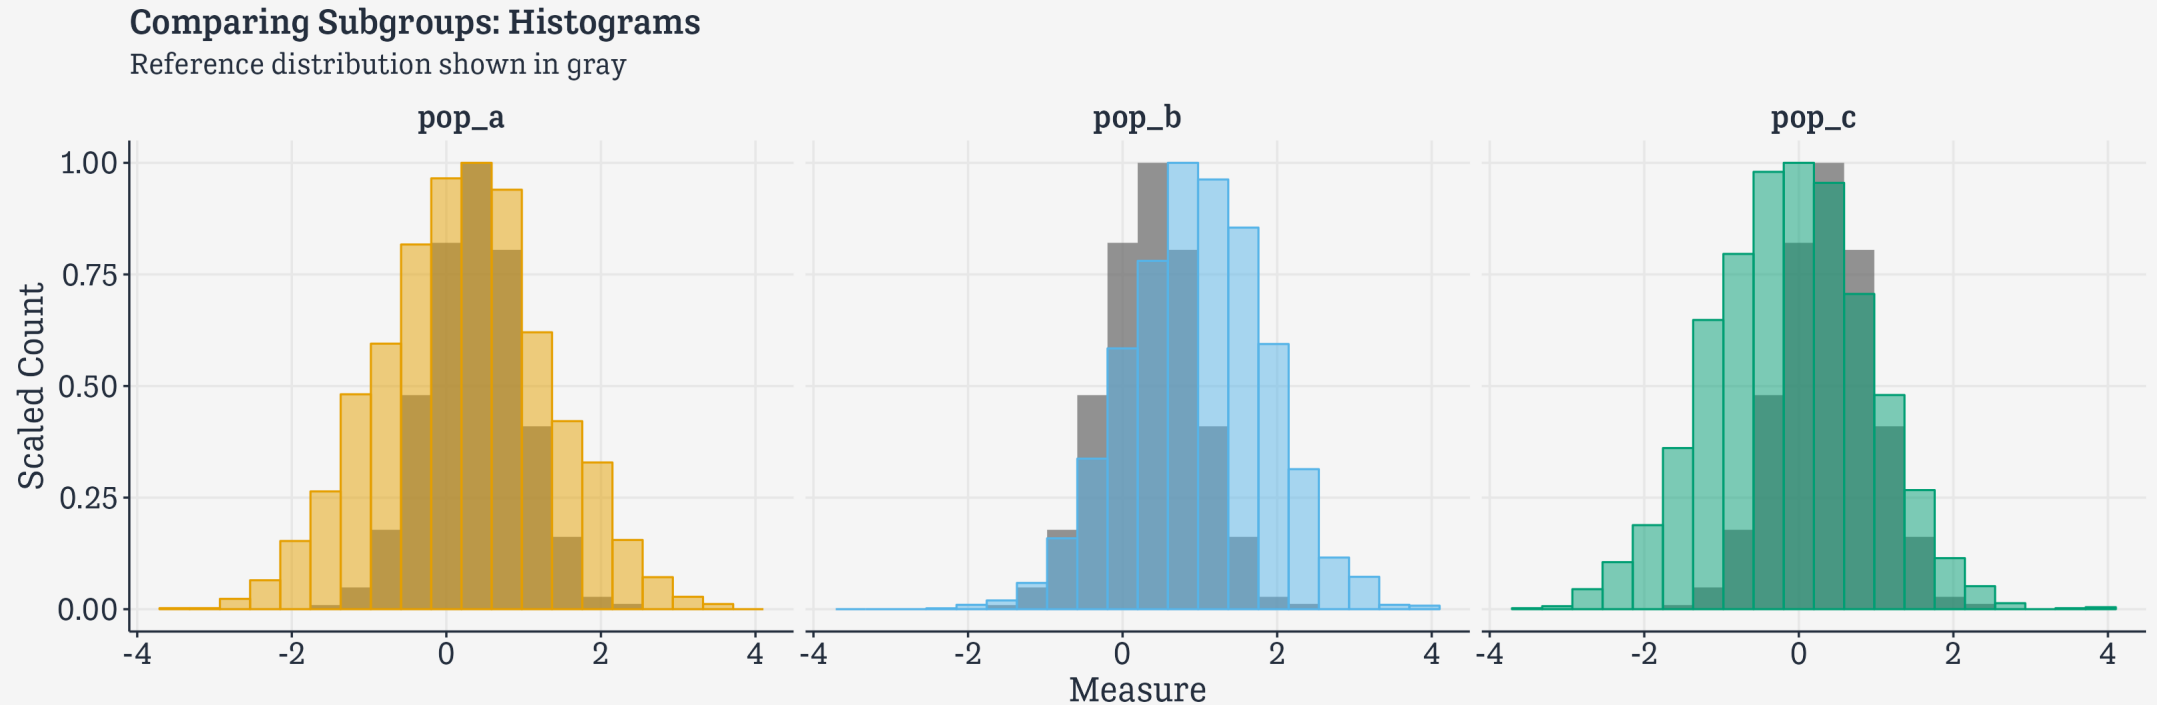
Now facet with that data

```
p_out <- df >
  pivot_longer(cols = pop_a:pop_c) >
  ggplot() +
  geom_histogram(mapping = aes(x = pop_total,
                              y = after_stat(ncount)),
                bins = 20, alpha = 0.7,
                fill = "gray40", size = 0.5) +
  geom_histogram(mapping = aes(x = value,
                              y = after_stat(ncount),
                              color = name, fill = name),
                stat = "bin", bins = 20, size = 0.5,
                alpha = 0.5) +
  guides(color = "none", fill = "none") +
  labs(x = "Measure", y = "Scaled Count",
       title = "Comparing Subgroups: Histograms",
       subtitle = "Reference distribution shown in gray") +
  facet_wrap(~ name, nrow = 1)
```

Remember, we can layer geoms one on top of the other. Here we call `geom_histogram()` twice. What happens if you comment one or other of them out?

The call to `guides()` turns off the legend for the color and fill, because we don't need them.

Now facet with that data



**Avoid counting up,
when necessary**

Sometimes no counting is required

```
titanic
```

```
##      fate    sex    n percent
## 1 perished  male 1364    62.0
## 2 perished female  126     5.7
## 3 survived  male  367    16.7
## 4 survived female 344    15.6
```

Here we just have a summary table and want to plot a few numbers directly in a bar chart.

geom_bar() wants to count up

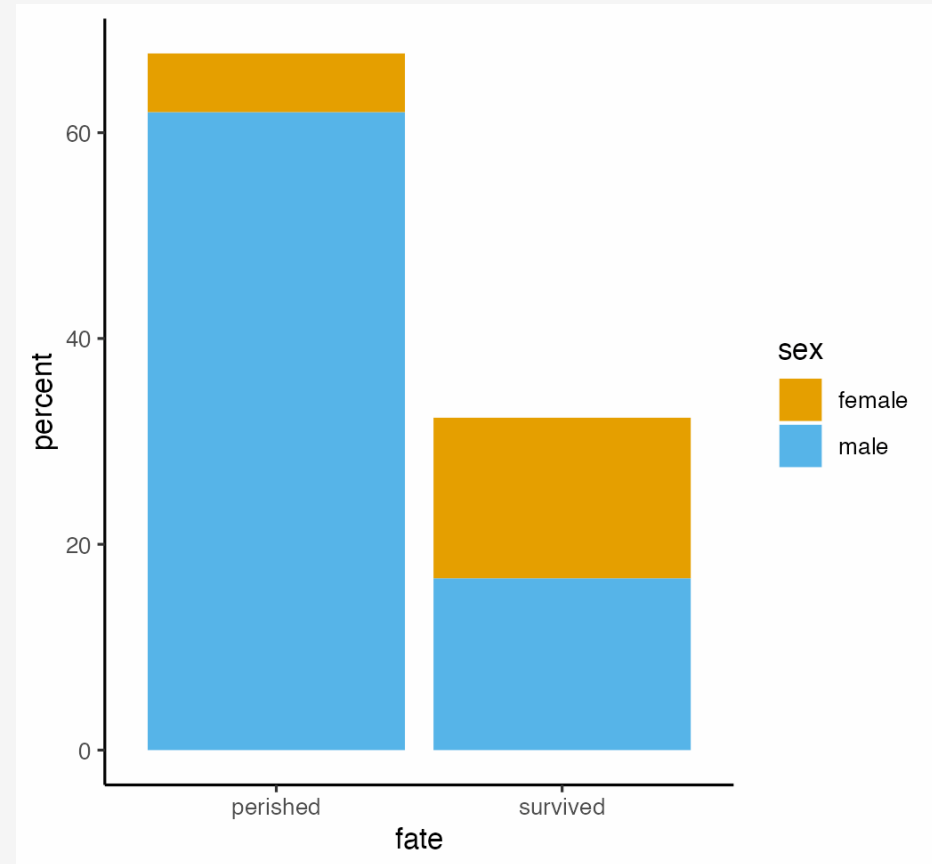
```
p <- ggplot(data = titanic,  
            mapping = aes(x = fate,  
                          y = percent,  
                          fill = sex))  
p + geom_bar(stat = "identity")
```

By default `geom_bar()` tries to count up data by category. (Really it's the `stat_count()` function that does this behind the scenes.) By saying `stat="identity"` we explicitly tell it not to do that. This also allows us to use a y mapping. Normally this would be the result of the counting up.

geom_bar() wants to count up

```
p <- ggplot(data = titanic,  
            mapping = aes(x = fate,  
                          y = percent,  
                          fill = sex))  
p + geom_bar(stat = "identity")
```

By default `geom_bar()` tries to count up data by category. (Really it's the `stat_count()` function that does this behind the scenes.) By saying `stat="identity"` we explicitly tell it not to do that. This also allows us to use a `y` mapping. Normally this would be the result of the counting up.



geom_bar() stacks bars by default

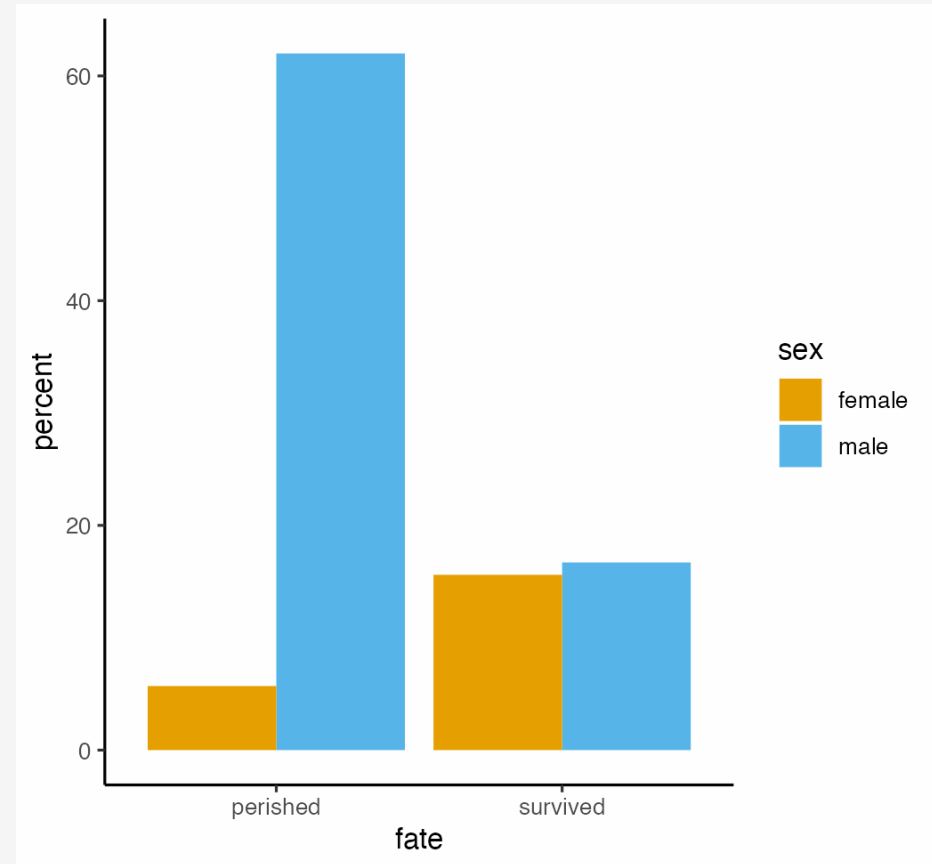
```
p <- ggplot(data = titanic,  
            mapping = aes(x = fate,  
                          y = percent,  
                          fill = sex))  
p + geom_bar(stat = "identity",  
            position = "dodge")
```

Position arguments adjust whether the things drawn are placed on top of one another ("stack"), side-by-side ("dodge"), or taken as-is ("identity").

geom_bar() stacks bars by default

```
p <- ggplot(data = titanic,  
            mapping = aes(x = fate,  
                          y = percent,  
                          fill = sex))  
p + geom_bar(stat = "identity",  
            position = "dodge")
```

Position arguments adjust whether the things drawn are placed on top of one another ("stack"), side-by-side ("dodge"), or taken as-is ("identity").



A quick **theme()** adjustment

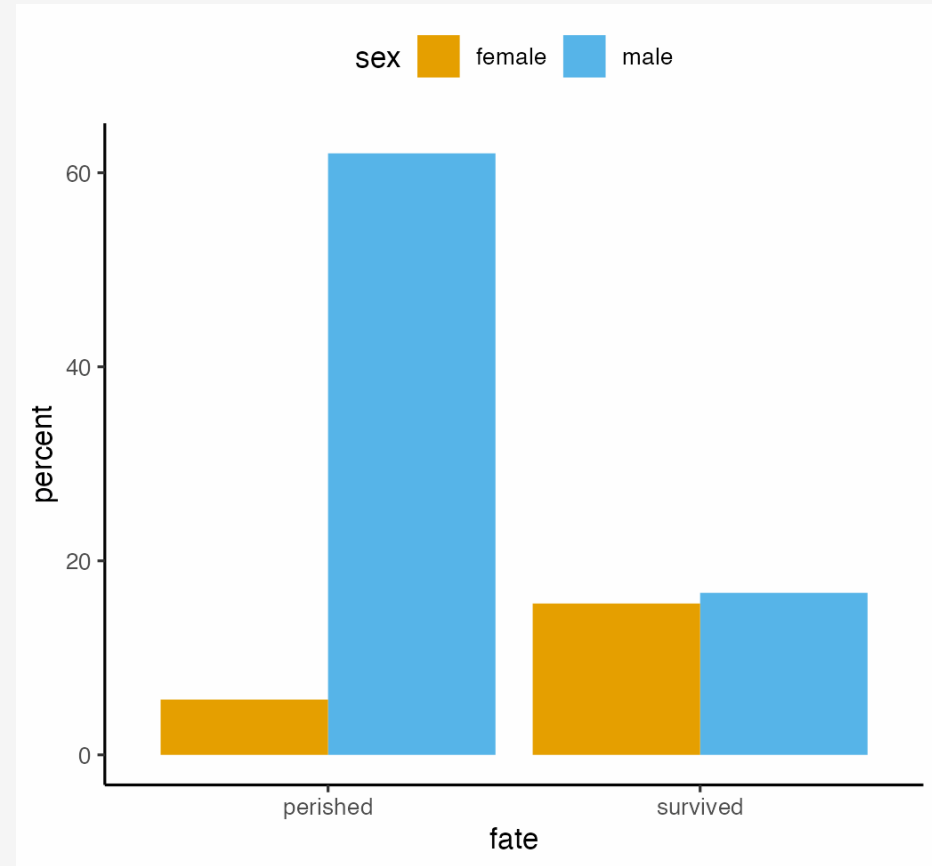
```
p <- ggplot(data = titanic,  
            mapping = aes(x = fate,  
                          y = percent,  
                          fill = sex))  
p + geom_bar(stat = "identity",  
            position = "dodge") +  
  theme(legend.position = "top")
```

The **theme()** function controls the styling of parts of the plot that don't belong to its "grammatical" structure. That is, that are not contributing to directly representing data.

A quick `theme()` adjustment

```
p <- ggplot(data = titanic,  
            mapping = aes(x = fate,  
                          y = percent,  
                          fill = sex))  
  
p + geom_bar(stat = "identity",  
            position = "dodge") +  
  theme(legend.position = "top")
```

The `theme()` function controls the styling of parts of the plot that don't belong to its "grammatical" structure. That is, that are not contributing to directly representing data.



For convenience, use `geom_col()`

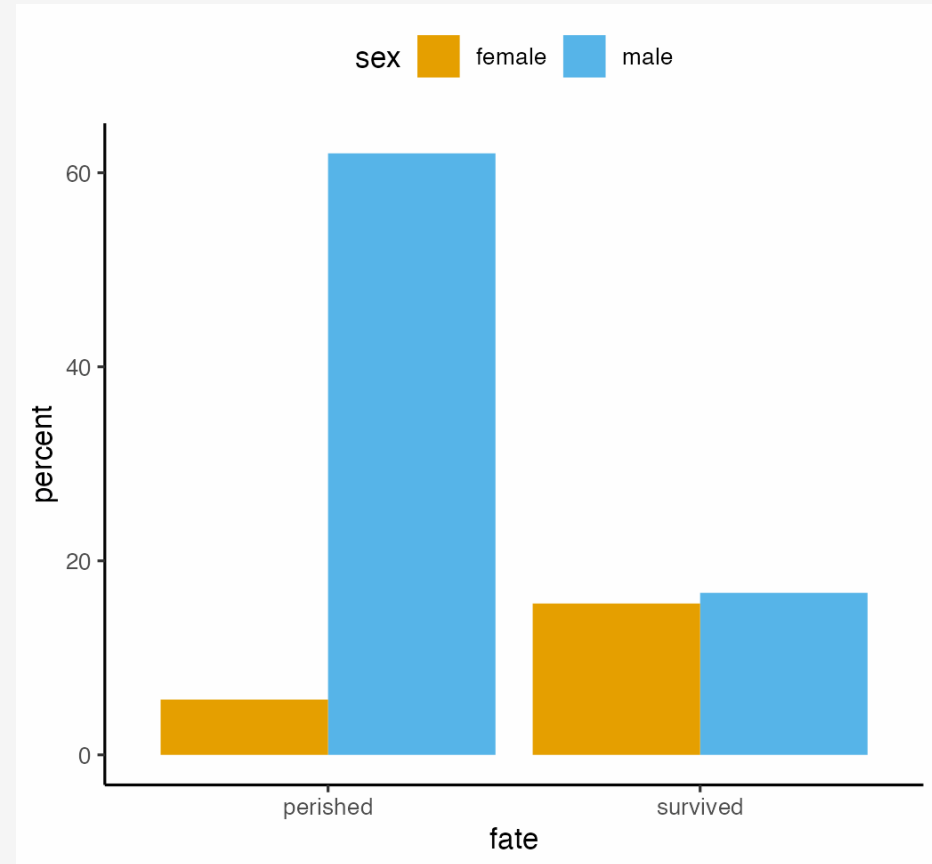
```
p <- ggplot(data = titanic,  
            mapping = aes(x = fate,  
                          y = percent,  
                          fill = sex))  
p + geom_col(position = "dodge") +  
  theme(legend.position = "top")
```

`geom_col()` assumes `stat = "identity"` by default. It's for when you want to directly plot a table of values, rather than create a bar chart by summing over one variable categorized by another.

For convenience, use `geom_col()`

```
p <- ggplot(data = titanic,  
            mapping = aes(x = fate,  
                          y = percent,  
                          fill = sex))  
p + geom_col(position = "dodge") +  
  theme(legend.position = "top")
```

`geom_col()` assumes `stat = "identity"` by default. It's for when you want to directly plot a table of values, rather than create a bar chart by summing over one variable categorized by another.



Using `geom_col()` for thresholds

```
oecd_sum
```

```
## # A tibble: 57 × 5
## # Groups:   year [57]
##   year other  usa  diff hi_lo
##   <int> <dbl> <dbl> <dbl> <chr>
## 1  1960  68.6  69.9  1.30 Below
## 2  1961  69.2  70.4  1.20 Below
## 3  1962  68.9  70.2  1.30 Below
## 4  1963  69.1  70    0.900 Below
## 5  1964  69.5  70.3  0.800 Below
## 6  1965  69.6  70.3  0.700 Below
## 7  1966  69.9  70.3  0.400 Below
## 8  1967  70.1  70.7  0.600 Below
## 9  1968  70.1  70.4  0.300 Below
## 10 1969  70.1  70.6  0.5    Below
## # i 47 more rows
```

Data comparing U.S. average life expectancy to the rest of the OECD average.

`diff` is difference in years with respect to the U.S.

`hi_lo` is a flag saying whether the OECD is above or below the U.S.

Using `geom_col()` for thresholds

```
p <- ggplot(data = oecd_sum,
            mapping = aes(x = year,
                          y = diff,
                          fill = hi_lo))

p_out <- p + geom_col() +
  geom_hline(yintercept = 0, linewidth = 1.2) +
  guides(fill = "none") +
  labs(x = NULL,
       y = "Difference in Years",
       title = "The U.S. Life Expectancy Gap",
       subtitle = "Difference between U.S. and
OECD average life expectancies, 1960-2015",
       caption = "Data: OECD.")
```

`geom_hline()` doesn't take any data argument. It just draws a horizontal line with a given y-intercept.

`x = NULL` means "Don't label the x-axis (not even with the default value, the variable name)."

Using `geom_col()` for thresholds

