

Show the **Right Numbers**

Data Visualization: Session 4

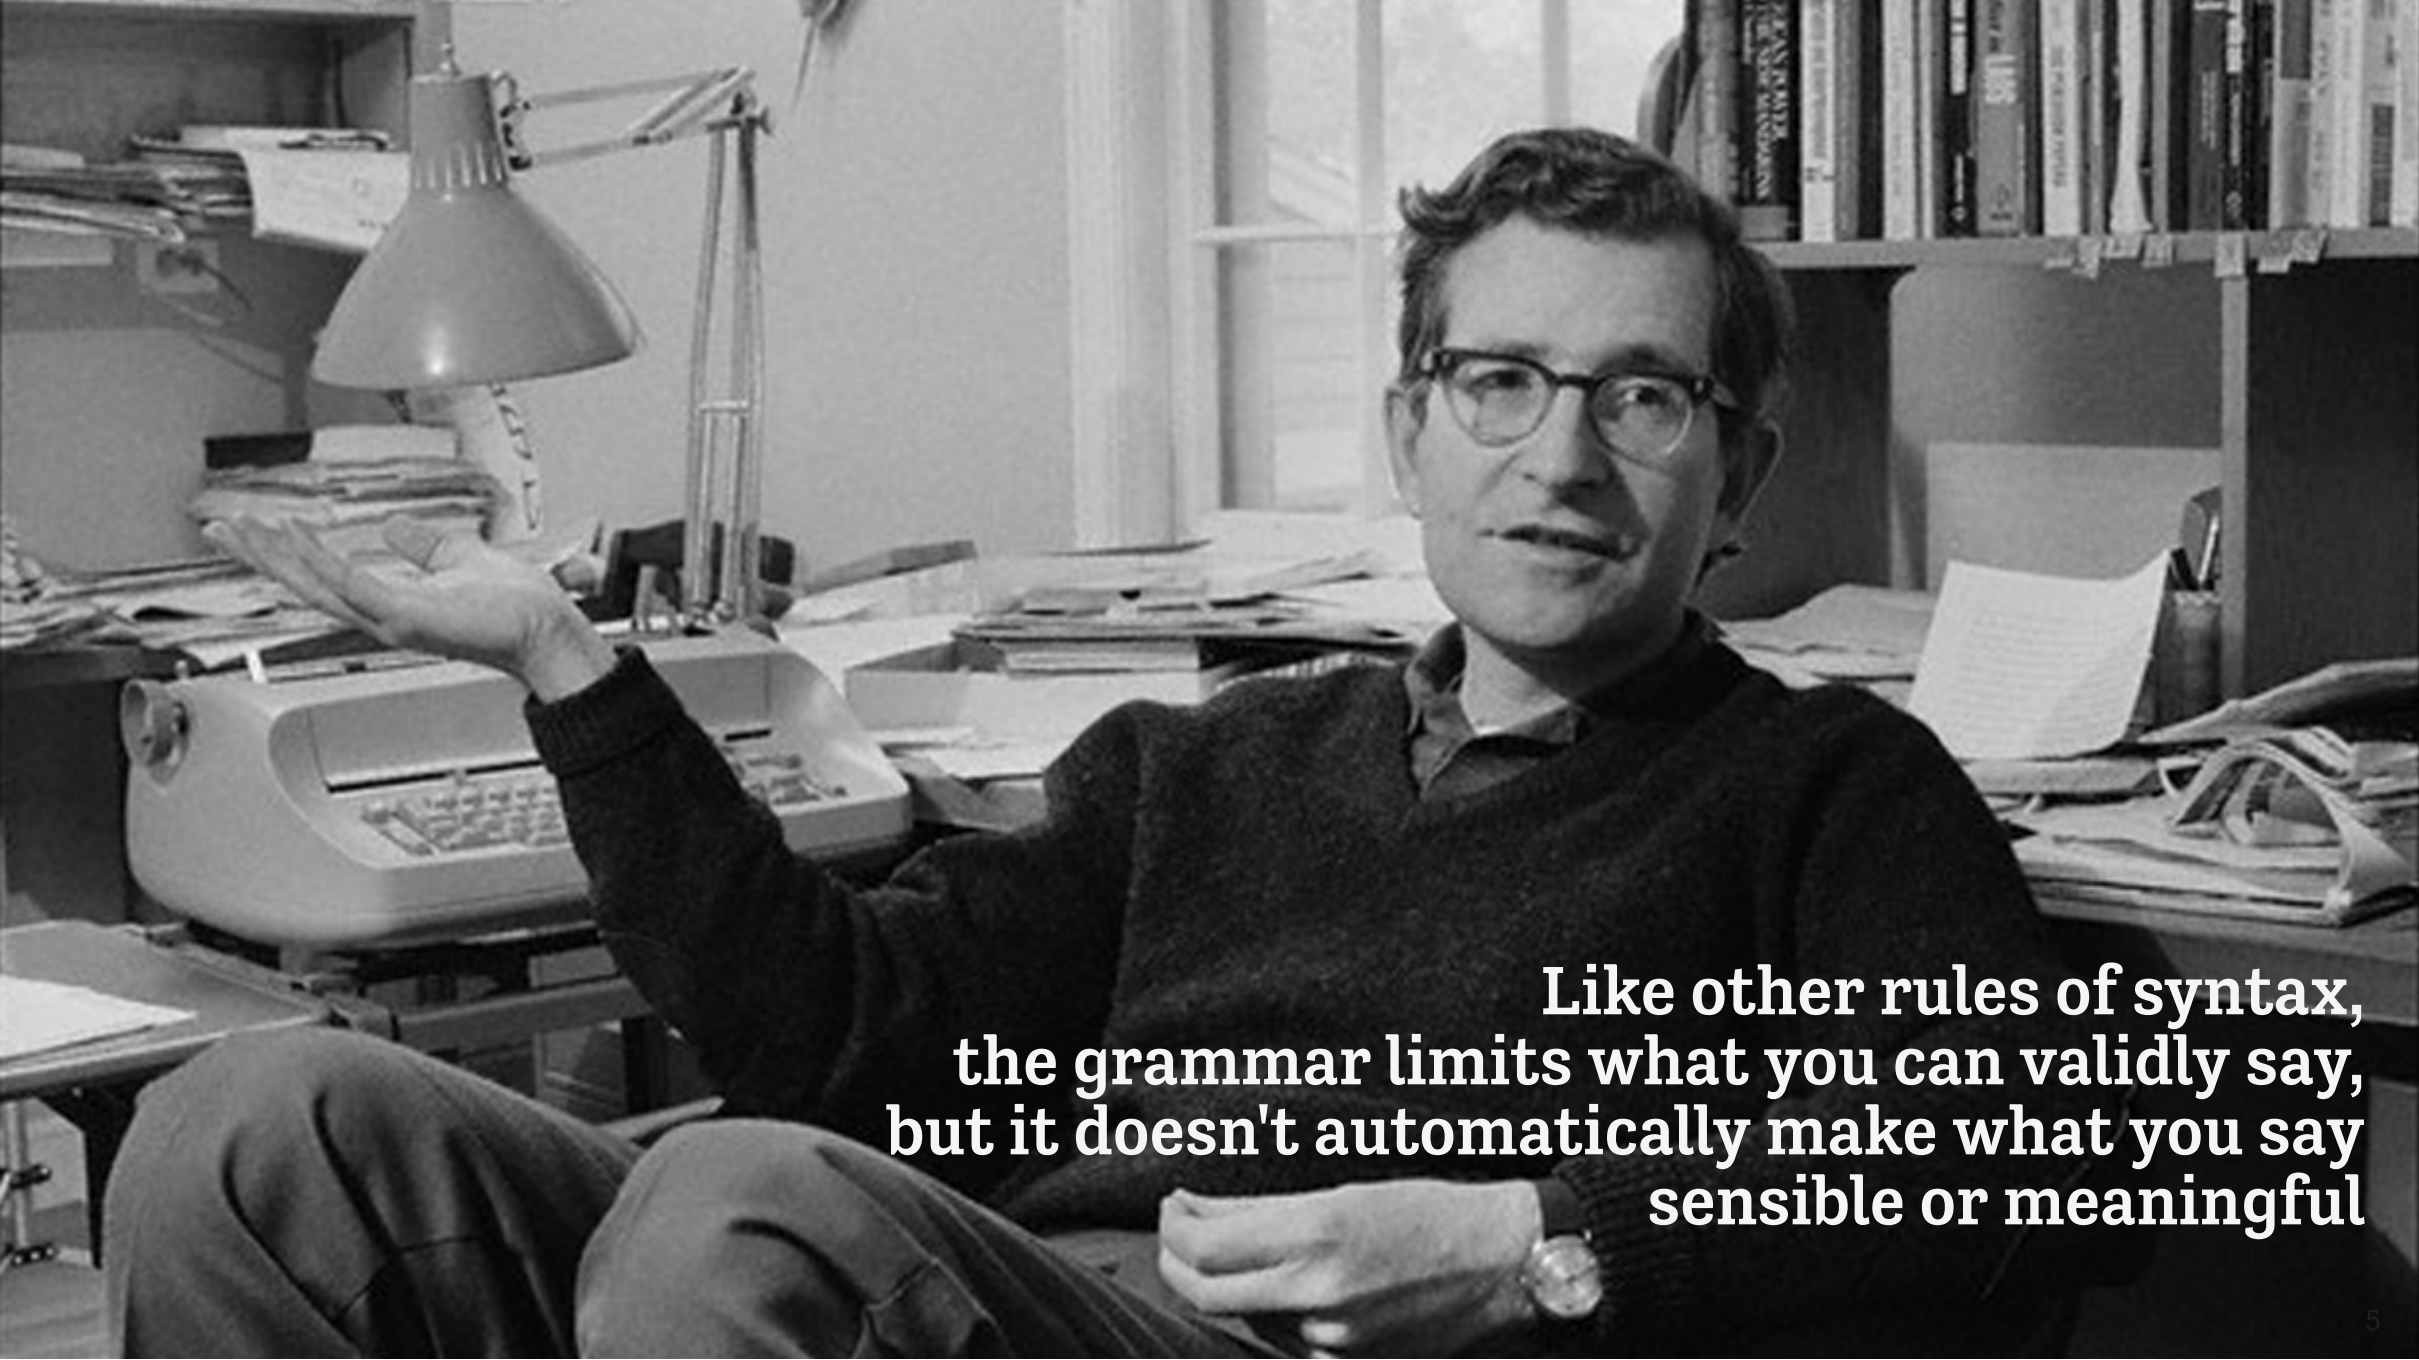
Kieran Healy
Code Horizons, May 2022

Set up our workspace

```
library(tidyverse)    # Your friend and mine  
library(gapminder)   # Gapminder data  
library(here)        # Portable file paths  
library(socviz)      # Handy socviz functions
```

ggplot
implements a
grammar of
graphics

The grammar is a set of rules for how to produce graphics from data, by *mapping* data to or *representing* it by geometric objects (like points and lines) that have aesthetic attributes (like position, color, size, and shape), together with further rules for transforming data if needed, for adjusting scales and their guides, and for projecting results onto some coordinate system.



**Like other rules of syntax,
the grammar limits what you can validly say,
but it doesn't automatically make what you say
sensible or meaningful**


Grouped data and the group aesthetic

Try to make a lineplot

```
p <- ggplot(data = gapminder,  
            mapping = aes(x = year,  
                          y = gdpPercap))
```

Try to make a lineplot

```
p <- ggplot(data = gapminder,  
            mapping = aes(x = year,  
                          y = gdpPercap))  
p + geom_line()
```




Try to make a lineplot

```
p <- ggplot(data = gapminder,  
            mapping = aes(x = year,  
                          y = gdpPercap))
```

Try to make a lineplot

```
p <- ggplot(data = gapminder,  
            mapping = aes(x = year,  
                          y = gdpPercap))  
  
## Geoms can take their own mappings, remember  
p + geom_line(mapping = aes(group = country))
```




Facet the plot

```
p <- ggplot(data = gapminder,  
            mapping = aes(x = year,  
                          y = gdpPercap))
```


Facet the plot

```
p <- ggplot(data = gapminder,  
            mapping = aes(x = year,  
                          y = gdpPercap))  
  
## Geoms can take their own mappings, remember  
p + geom_line(mapping = aes(group = country))
```



Facet the plot

```
p <- ggplot(data = gapminder,  
            mapping = aes(x = year,  
                          y = gdpPercap))  
  
## Geoms can take their own mappings, remember  
p + geom_line(mapping = aes(group = country)) +  
  facet_wrap(~ continent)
```



Faceting is a very powerful tool

A facet is not a geom; it's a way of arranging repeated geoms by some additional variable

Faceting is a very powerful tool

A facet is not a geom; it's a way of arranging repeated geoms by some additional variable

Facets use R's "formula" syntax: `facet_wrap(~ continent)`

Faceting is a very powerful tool

A facet is not a geom; it's a way of arranging repeated geoms by some additional variable

Facets use R's "formula" syntax: `facet_wrap(~ continent)`

Read the `~` as "on" or "by"

Faceting is a very powerful tool

A facet is not a geom; it's a way of arranging repeated geoms by some additional variable

Facets use R's "formula" syntax: `facet_wrap(~ continent)`

Read the `~` as "on" or "by"

You can also use this syntax: `facet_wrap(vars(continent))`

Faceting is a very powerful tool

A facet is not a geom; it's a way of arranging repeated geoms by some additional variable

Facets use R's "formula" syntax: `facet_wrap(~ continent)`

Read the `~` as "on" or "by"


You can also use this syntax: `facet_wrap(vars(continent))`

This is newer, and consistent with other ways of referring to variables within tidyverse functions.

Facets in action

```
p <- ggplot(data = gapminder,  
            mapping = aes(x = year,  
                          y = gdpPercap))  
  
p_out <- p + geom_line(color="gray70",  
                      mapping=aes(group = country)) +  
  geom_smooth(size = 1.1,  
             method = "loess",  
             se = FALSE) +  
  scale_y_log10(labels=scales::label_dollar()) +  
  facet_wrap(~ continent, ncol = 5) +  
  labs(x = "Year",  
       y = "log GDP per capita",  
       title = "GDP per capita on Five Continents",  
       caption = "Data: Gapminder")
```

GDP per capita on Five Continents



Data: Gapminder

A more polished faceted plot.

One-variable summaries

midwest

County-level census data for Midwestern U.S. Counties

```
midwest
```

```
## # A tibble: 437 × 28
```

```
##       PID county    state  area poptotal popdensity popwhite popblack popamerindian popasian popother percwhite percblack percame  
##    <int> <chr>    <chr> <dbl>    <int>    <dbl>    <int>    <int>          <int>    <int>    <int>    <dbl>    <dbl>  
##  1   561 ADAMS      IL    0.052    66090    1271.    63917    1702            98      249     124     96.7     2.58  
##  2   562 ALEXANDER IL    0.014    10626     759     7054    3496            19      48      9     66.4    32.9  
##  3   563 BOND       IL    0.022    14991     681.    14477     429            35      16     34     96.6     2.86  
##  4   564 BOONE      IL    0.017    30806    1812.    29344     127            46     150    1139     95.3     0.412  
##  5   565 BROWN      IL    0.018     5836     324.    5264     547            14       5      6     90.2     9.37  
##  6   566 BUREAU     IL    0.05    35688     714.    35157      50            65     195    221     98.5     0.140  
##  7   567 CALHOUN    IL    0.017     5322     313.    5298       1            8      15      0     99.5     0.0188  
##  8   568 CARROLL    IL    0.027    16805     622.    16519     111            30      61     84     98.3     0.661  
##  9   569 CASS       IL    0.024    13437     560.    13384      16            8      23      6     99.6     0.119  
## 10  570 CHAMPAIGN IL    0.058   173025    2983.   146506   16559           331    8033   1596     84.7     9.57  
## # ... with 427 more rows, and 13 more variables: percother <dbl>, popadults <int>, perchs <dbl>, percollege <dbl>, percprof <dbl>  
## #   poppovertyknown <int>, percpovertyknown <dbl>, percbelowpoverty <dbl>, percchildbelowpovert <dbl>, percadultpoverty <dbl>,  
## #   percelderlypoverty <dbl>, inmetro <int>, category <chr>
```

stat_ functions work behind the scenes

```
p <- ggplot(data = midwest,  
            mapping = aes(x = area))
```

```
p + geom_histogram()
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```

Here the default `stat_` function for this geom has to make a choice. It is letting us know we might want to override it.


stat_ functions work behind the scenes

```
p <- ggplot(data = midwest,  
            mapping = aes(x = area))
```

```
p + geom_histogram()
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```

Here the default `stat_` function for this geom has to make a choice. It is letting us know we might want to override it.



stat_ functions work behind the scenes


```
p <- ggplot(data = midwest,  
            mapping = aes(x = area))  
  
p + geom_histogram(bins = 10)
```

We can choose *either* the number of bins *or* the binwidth

stat_ functions work behind the scenes

```
p <- ggplot(data = midwest,  
            mapping = aes(x = area))  
  
p + geom_histogram(bins = 10)
```

We can choose *either* the number of bins *or* the binwidth



Compare two distributions

```
## Two state codes
oh_wi <- c("OH", "WI")

midwest |>
  filter(state %in% oh_wi) |>
  ggplot(mapping = aes(x = percollege,
                       fill = state)) +
  geom_histogram(alpha = 0.5,
                 position = "identity")
```

Here we do the whole thing in a **pipeline** using the pipe and the dplyr verb **filter()** to subset rows of the data by some condition.


Experiment with leaving the `position` argument out, or changing it to "dodge".

Compare two distributions

```
## Two state codes
oh_wi <- c("OH", "WI")

midwest |>
  filter(state %in% oh_wi) |>
  ggplot(mapping = aes(x = percollege,
                        fill = state)) +
  geom_histogram(alpha = 0.5,
                 position = "identity")
```

Here we do the whole thing in a **pipeline** using the pipe and the dplyr verb **filter()** to subset rows of the data by some condition. Experiment with leaving the **position** argument out, or changing it to "dodge".




geom_hist()'s counterpart, geom_density()

```
p <- ggplot(data = midwest,  
            mapping = aes(x = area))  
  
p + geom_density()
```

geom_hist()'s counterpart, **geom_density()**

```
p <- ggplot(data = midwest,  
            mapping = aes(x = area))  
  
p + geom_density()
```




geom_hist()'s counterpart, geom_density()

```
p <- ggplot(data = midwest,  
            mapping = aes(x = area,  
                          fill = state,  
                          color = state))  
p + geom_density(alpha = 0.3)
```

geom_hist()'s counterpart, geom_density()

```
p <- ggplot(data = midwest,  
  mapping = aes(x = area,  
    fill = state,  
    color = state))  
p + geom_density(alpha = 0.3)
```



geom_hist()'s counterpart, geom_density()


```
midwest |>
  filter(state %in% oh_wi) |>
  ggplot(mapping = aes(x = area,
                       fill = state,
                       color = state)) +
  geom_density(mapping = aes(y = ..ndensity..),
              alpha = 0.4)
```

`..ndensity..` here is not in our data! It's *computed*. Histogram and density geoms have default statistics, but you can ask them to do more. The `stat_` functions associated with each `geom_` do this work behind the scenes.

geom_hist()'s counterpart, geom_density()

```
midwest |>
  filter(state %in% oh_wi) |>
  ggplot(mapping = aes(x = area,
                        fill = state,
                        color = state)) +
  geom_density(mapping = aes(y = ..ndensity..),
               alpha = 0.4)
```

`..ndensity..` here is not in our data! It's *computed*. Histogram and density geoms have default statistics, but you can ask them to do more. The `stat_` functions associated with each `geom_` do this work behind the scenes.



**Compare subgroups to a
reference distribution**

Some made-up data

Consider 3,000 observations of some unit (e.g., a county) with summary measures for each group, and the population average.

df

```
## # A tibble: 3,000 × 5
##   unit  pop_a pop_b  pop_c pop_total
##   <int> <dbl> <dbl>   <dbl>   <dbl>
## 1     1  1.29  1.93  -0.0869  1.09
## 2     2  0.522  0.536 -0.762   0.190
## 3     3  2.14   1.47  -0.616   1.15
## 4     4  1.13   0.673 -0.242   0.575
## 5     5  1.04   1.30   1.18    1.12
## 6     6  1.80   0.140  2.05    1.33
## 7     7  0.186  1.30  -0.709   0.476
## 8     8 -0.953  0.520  -2.44   -0.767
## 9     9  0.700  1.66  -1.09    0.749
## 10    10  0.0416 0.484 -0.180   0.177
## # ... with 2,990 more rows
```

First effort: Hard to read


```
df |>
  pivot_longer(cols = pop_a:pop_total) |>
  ggplot() +
  geom_histogram(mapping = aes(x = value,
                              y = ..ncount..,
                              color = name, fill = name),
                 stat = "bin", bins = 20, size = 0.5,
                 alpha = 0.7,
                 position = "identity") +
  labs(x = "Measure", y = "Scaled Count", color = "Group",
       fill = "Group",
       title = "Comparing Subgroups: Histograms")
```

Again, `..ncount..` is computed. The periods on either side are just a naming convention to show that the measure is computed by the `stat_` function (and to make sure it doesn't clash with any actual names in your data.)

First effort: Hard to read

```
df |>
  pivot_longer(cols = pop_a:pop_total) |>
  ggplot() +
  geom_histogram(mapping = aes(x = value,
                              y = ..ncount..,
                              color = name, fill = name),
                 stat = "bin", bins = 20, size = 0.5,
                 alpha = 0.7,
                 position = "identity") +
  labs(x = "Measure", y = "Scaled Count", color = "Group",
       fill = "Group",
       title = "Comparing Subgroups: Histograms")
```

Again, `..ncount..` is computed. The periods on either side are just a naming convention to show that the measure is computed by the `stat_` function (and to make sure it doesn't clash with any actual names in your data.)




Try faceting instead

```
p_out <- df |>
  pivot_longer(cols = pop_a:pop_c) |>
  ggplot() +
  geom_histogram(mapping = aes(x = pop_total,
                              y = ..ncount..),
                bins = 20, alpha = 0.7,
                fill = "gray40", size = 0.5) +
  geom_histogram(mapping = aes(x = value,
                              y = ..ncount..,
                              color = name, fill = name),
                stat = "bin", bins = 20, size = 0.5,
                alpha = 0.5) +
  guides(color = "none", fill = "none") +
  labs(x = "Measure", y = "Scaled Count",
       title = "Comparing Subgroups: Histograms",
       subtitle = "Reference distribution shown in gray") +
  facet_wrap(~ name, nrow = 1)
```

Something we haven't seen before, but will be using a lot: We can layer geoms one on top of the other. Here we call `geom_histogram()` twice. What happens if you comment one or other of them out?

The call to `guides()` turns off the legend for the color and fill, because we don't need them.

Try faceting instead



**Avoid counting up, if
necessary**

Sometimes no counting is required

```
titanic
```

```
##      fate    sex    n percent
## 1 perished  male 1364    62.0
## 2 perished  female 126     5.7
## 3 survived  male  367    16.7
## 4 survived  female 344    15.6
```

Here we just have a summary table and want to plot a few numbers directly in a bar chart.

geom_bar() wants to count up


```
p <- ggplot(data = titanic,  
            mapping = aes(x = fate,  
                          y = percent,  
                          fill = sex))  
p + geom_bar(stat = "identity")
```

By default `geom_bar()` tries to count up data by category. By saying `stat="identity"` we explicitly tell it not to do that. This also allows us to use a `y` mapping, because normally this would be determined by the counting up.

geom_bar() wants to count up

```
p <- ggplot(data = titanic,  
            mapping = aes(x = fate,  
                          y = percent,  
                          fill = sex))  
p + geom_bar(stat = "identity")
```

By default `geom_bar()` tries to count up data by category. By saying `stat="identity"` we explicitly tell it not to do that. This also allows us to use a `y` mapping, because normally this would be determined by the counting up.



geom_bar() stacks bars by default


```
p <- ggplot(data = titanic,  
            mapping = aes(x = fate,  
                          y = percent,  
                          fill = sex))  
p + geom_bar(stat = "identity",  
            position = "dodge")
```

Position arguments adjust whether the things drawn are placed on top of one another ("stack"), side-by-side ("dodge"), or taken as-is ("identity").

geom_bar() stacks bars by default

```
p <- ggplot(data = titanic,  
            mapping = aes(x = fate,  
                          y = percent,  
                          fill = sex))  
p + geom_bar(stat = "identity",  
            position = "dodge")
```

Position arguments adjust whether the things drawn are placed on top of one another ("stack"), side-by-side ("dodge"), or taken as-is ("identity").



A quick **theme()** adjustment


```
p <- ggplot(data = titanic,  
            mapping = aes(x = fate,  
                          y = percent,  
                          fill = sex))  
  
p + geom_bar(stat = "identity",  
            position = "dodge") +  
  theme(legend.position = "top")
```

The **theme()** function controls the styling of parts of the plot that don't belong to its "grammatical" structure. That is, that are not contributing to directly representing data.

A quick `theme()` adjustment

```
p <- ggplot(data = titanic,  
  mapping = aes(x = fate,  
    y = percent,  
    fill = sex))  
p + geom_bar(stat = "identity",  
  position = "dodge") +  
  theme(legend.position = "top")
```

The `theme()` function controls the styling of parts of the plot that don't belong to its "grammatical" structure. That is, that are not contributing to directly representing data.



For convenience, use `geom_col()`


```
p <- ggplot(data = titanic,  
            mapping = aes(x = fate,  
                          y = percent,  
                          fill = sex))  
p + geom_col(position = "dodge") +  
  theme(legend.position = "top")
```

`geom_col()` assumes `stat = "identity"` by default. It's for when you want to directly plot a table of values.

For convenience, use `geom_col()`

```
p <- ggplot(data = titanic,  
            mapping = aes(x = fate,  
                          y = percent,  
                          fill = sex))  
p + geom_col(position = "dodge") +  
  theme(legend.position = "top")
```

`geom_col()` assumes `stat = "identity"` by default. It's for when you want to directly plot a table of values.



Using `geom_col()` for thresholds

```
oecd_sum
```

```
## # A tibble: 57 × 5
## # Groups:   year [57]
##   year other   usa  diff hi_lo
##   <int> <dbl> <dbl> <dbl> <chr>
## 1  1960  68.6  69.9  1.30 Below
## 2  1961  69.2  70.4  1.20 Below
## 3  1962  68.9  70.2  1.30 Below
## 4  1963  69.1  70    0.900 Below
## 5  1964  69.5  70.3  0.800 Below
## 6  1965  69.6  70.3  0.700 Below
## 7  1966  69.9  70.3  0.400 Below
## 8  1967  70.1  70.7  0.600 Below
## 9  1968  70.1  70.4  0.300 Below
## 10 1969  70.1  70.6  0.5    Below
## # ... with 47 more rows
```

Data comparing U.S. average life expectancy to the rest of the OECD average.

`diff` is difference in years with respect to the U.S.

`hi_lo` is a flag saying whether the OECD is above or below the U.S.

Using `geom_col()` for thresholds

```
p <- ggplot(data = oecd_sum,
            mapping = aes(x = year,
                          y = diff,
                          fill = hi_lo))

p_out <- p + geom_col() +
  geom_hline(yintercept = 0, size = 1.2) +
  guides(fill = "none") +
  labs(x = NULL,
       y = "Difference in Years",
       title = "The U.S. Life Expectancy Gap",
       subtitle = "Difference between U.S. and
OECD average life expectancies, 1960-2015",
       caption = "Data: OECD.")
```

`geom_hline()` draws a horizontal line with a given y-intercept.
x = NULL means "Don't label the x-axis (not even with the default value, the variable name)."

Using `geom_col()` for thresholds

