

Data Visualization - 6.

Work With Models

Kieran Healy
Code Horizons

February 2025

Work with Models

Load our packages

```
library(here)      # manage file paths
library(socviz)    # data and some useful functions
library(tidyverse) # your friend and mine
library(gapminder) # Everyone's favorite dataset
library(broom)     # Tidy model output
library(marginaleffects) # Tidy marginal effects
library(modelsummary) # Tidy summary tables and graphs
library(scales)    # Format our axes and guides
```

```
Attaching package: 'scales'
```

```
The following object is masked from 'package:purrr':
```

```
discard
```

```
The following object is masked from 'package:readr':
```


```
col_factor
```

ggplot can work with models

We know because `geoms` often do calculations in the background, via their `stat` functions.


```
p <- gapminder %>
  ggplot(mapping = aes(x = log(gdpPercap),
                        y = lifeExp))

p + geom_point(alpha=0.1) +
  geom_smooth(color = "tomato",
              fill="tomato",
              method = MASS::rlm) +
  geom_smooth(color = "steelblue",
              fill="steelblue",
              method = "lm") +
  labs(title = "Robust and OLS fits")
```




These can be complex ...

```
p + geom_point(alpha=0.1) +  
  geom_smooth(color = "tomato",  
  method = "lm",  
  size = 1.2,  
  formula = y ~ splines::bs(x, 3),  
  se = FALSE)
```



... but we usually won't do this in **ggplot**

```
p + geom_point(alpha=0.1) +  
  geom_quantile(color = "tomato",  
    size = 1.2,  
    method = "rqss",  
    lambda = 1,  
    quantiles = c(0.20, 0.5, 0.85))
```



**Transform and
summarize first.**

**Then send your
clean tables to
ggplot.**

Look inside the box

Objects are To-Do List Bento Boxes

```
gapminder
```

```
# A tibble: 1,704 × 6
  country   continent   year lifeExp     pop gdpPercap
  <fct>     <fct>     <int>   <dbl>   <int>     <dbl>
1 Afghanistan Asia      1952    28.8  8425333    779.
2 Afghanistan Asia      1957    30.3  9240934    821.
3 Afghanistan Asia      1962    32.0 10267083    853.
4 Afghanistan Asia      1967    34.0 11537966    836.
5 Afghanistan Asia      1972    36.1 13079460    740.
6 Afghanistan Asia      1977    38.4 14880372    786.
7 Afghanistan Asia      1982    39.9 12881816    978.
8 Afghanistan Asia      1987    40.8 13867957    852.
9 Afghanistan Asia      1992    41.7 16317921    649.
10 Afghanistan Asia     1997    41.8 22227415    635.
# i 1,694 more rows
```

Fit a model

```
out ← lm(formula = lifeExp ~ gdpPercap + log(pop) + continent,  
         data = gapminder)  
  
summary(out)
```

Call:

```
lm(formula = lifeExp ~ gdpPercap + log(pop) + continent, data = gapminder)
```

Residuals:

Min	1Q	Median	3Q	Max
-47.490	-4.614	0.250	5.293	26.094

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	3.816e+01	2.050e+00	18.618	< 2e-16 **
gdpPercap	4.557e-04	2.345e-05	19.435	< 2e-16 **
log(pop)	6.394e-01	1.329e-01	4.810	1.64e-06 **
continentAmericas	1.308e+01	6.063e-01	21.579	< 2e-16 **
continentAsia	7.784e+00	5.810e-01	13.398	< 2e-16 **
continentEurope	1.695e+01	6.350e-01	26.691	< 2e-16 **
continentOceania	1.764e+01	1.779e+00	9.916	< 2e-16 **

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Poke around inside

```
out
  └── coefficients
  └── residuals
  └── effects
  └── rank
    └── qr
      └── qr
      └── pivot
      └── qraux
      └── tol
      └── rank
  └── df.residual
  └── contrasts
  └── xlevels
  └── call
  └── terms
  └── model.frame
```

Use the Object Inspector to take a look

Predict from models: DIY method

Behind the curtain

`predict()` and its methods do a lot of work behind the scenes

We won't usually need to do this stuff manually. But the idea is that the generic `predict()` function has *methods* for specific sorts of models. Give it a model and some new data and it will produce predicted values for the new data. Here's an example.

The labor-intensive way

```
min_gdp ← min(gapminder$gdpPercap)
max_gdp ← max(gapminder$gdpPercap)
med_pop ← median(gapminder$pop)

# Make a grid of predictor values
pred_df ← expand_grid(gdpPercap = (seq(from = min_gdp,
                                         to = max_gdp,
                                         length.out = 100)),
                       pop = med_pop,
                       continent = c("Africa", "Americas",
                                     "Asia", "Europe", "Oceania"))

pred_df
```

```
# A tibble: 500 × 3
  gdpPercap      pop continent
  <dbl>     <dbl>   <chr>
1 241.    7023596. Africa
2 241.    7023596. Americas
3 241.    7023596. Asia
4 241.    7023596. Europe
5 241.    7023596. Oceania
6 1385.   7023596. Africa
7 1385.   7023596. Americas
8 1385.   7023596. Asia
9 1385.   7023596. Europe
10 1385.   7023596. Oceania
# i 490 more rows
```

The labor-intensive way

```
# Get the predicted values
pred_out ← predict(object = out,
                     newdata = pred_df,
                     interval = "confidence")
head(pred_out)

  fit      lwr      upr
1 48.35388 47.67735 49.03041
2 61.43646 60.43917 62.43375
3 56.13821 55.22045 57.05597
4 65.30361 64.21794 66.38927
5 65.99517 62.55277 69.43757
6 48.87530 48.20261 49.54799
```

The labor-intensive way

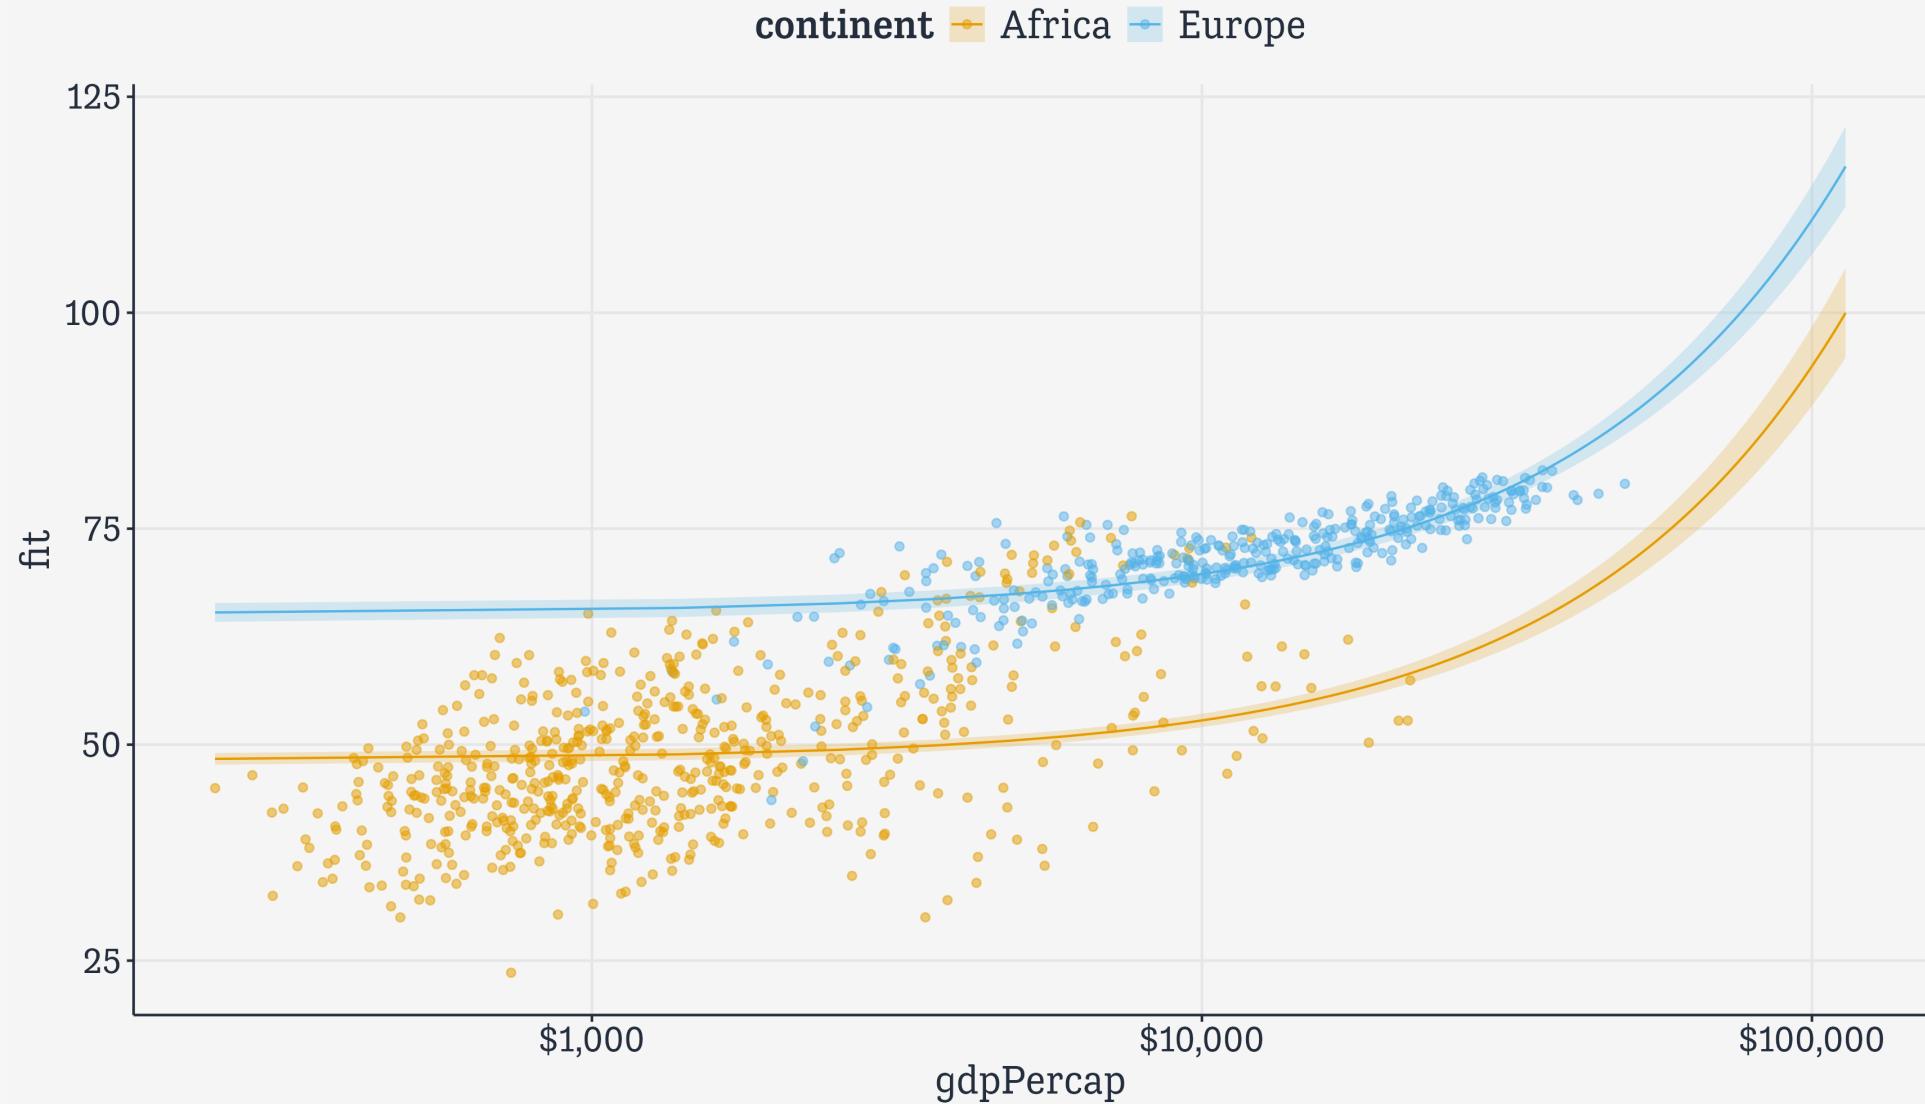
```
# Bind them into one data frame. We can do this safely  
# here because we know the row order by construction.  
# But this is not a safe approach in general.
```

```
pred_df ← cbind(pred_df, pred_out)  
head(pred_df)
```

	gdpPercap	pop	continent	fit	lwr	upr
1	241.1659	7023596	Africa	48.35388	47.67735	49.03041
2	241.1659	7023596	Americas	61.43646	60.43917	62.43375
3	241.1659	7023596	Asia	56.13821	55.22045	57.05597
4	241.1659	7023596	Europe	65.30361	64.21794	66.38927
5	241.1659	7023596	Oceania	65.99517	62.55277	69.43757
6	1385.4282	7023596	Africa	48.87530	48.20261	49.54799

The labor-intensive way

```
p ← ggplot(data = subset(pred_df, continent %in% c("Europe", "Africa")),  
            aes(x = gdpPercap,  
                 y = fit,  
                 ymin = lwr,  
                 ymax = upr,  
                 color = continent,  
                 fill = continent,  
                 group = continent))  
  
# Use the original data as the point layer  
p_out ← p + geom_point(data = subset(gapminder,  
                                         continent %in% c("Europe", "Africa")),  
                         mapping = aes(x = gdpPercap, y = lifeExp,  
                                        color = continent),  
                         alpha = 0.5,  
                         inherit.aes = FALSE) +  
# And the predicted values to draw the lines  
  geom_line() +  
  geom_ribbon(alpha = 0.2, color = FALSE) +  
  scale_x_log10(labels = scales::label_dollar())
```



Use broom to tidy models

We can't do anything with this

```
out ← lm(formula = lifeExp ~ gdpPercap + log(pop) + continent,  
         data = gapminder)  
  
summary(out)
```

Call:

```
lm(formula = lifeExp ~ gdpPercap + log(pop) + continent, data = gapminder)
```

Residuals:

Min	1Q	Median	3Q	Max
-47.490	-4.614	0.250	5.293	26.094

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	3.816e+01	2.050e+00	18.618	< 2e-16 **
gdpPercap	4.557e-04	2.345e-05	19.435	< 2e-16 **
log(pop)	6.394e-01	1.329e-01	4.810	1.64e-06 **
continentAmericas	1.308e+01	6.063e-01	21.579	< 2e-16 **
continentAsia	7.784e+00	5.810e-01	13.398	< 2e-16 **
continentEurope	1.695e+01	6.350e-01	26.691	< 2e-16 **
continentOceania	1.764e+01	1.779e+00	9.916	< 2e-16 **

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Tidy regression output with **broom**

```
library(broom)  
  
tidy(out)  
  
# A tibble: 7 × 5  
  term      estimate std.error statistic p.value  
  <chr>     <dbl>     <dbl>     <dbl>    <dbl>  
1 (Intercept) 38.2      2.05      18.6  1.50e- 70  
2 gdpPerCap   0.000456  0.0000234  19.4  3.98e- 76  
3 log(pop)    0.639     0.133      4.81  1.64e-  6  
4 continentAmericas 13.1     0.606      21.6  1.85e- 91  
5 continentAsia    7.78     0.581      13.4  5.52e- 39  
6 continentEurope   16.9     0.635      26.7  2.43e-131  
7 continentOceania  17.6     1.78       9.92  1.43e- 22
```

That's a *lot* nicer. Now it's just a tibble. We know those.

Tidy regression output with **broom**

```
out_conf ← tidy(out, conf.int = TRUE)  
out_conf
```

```
# A tibble: 7 × 7  
  term          estimate std.error statistic   p.value conf.low conf.high  
  <chr>        <dbl>     <dbl>      <dbl>    <dbl>     <dbl>      <dbl>  
1 (Intercept)  38.2      2.05       18.6  1.50e- 70  34.1      42.2  
2 gdpPercap    0.000456  0.0000234   19.4  3.98e- 76  0.000410  0.000502  
3 log(pop)     0.639      0.133       4.81  1.64e- 6   0.379      0.900  
4 continentAmericas 13.1     0.606       21.6  1.85e- 91  11.9      14.3  
5 continentAsia    7.78     0.581       13.4  5.52e- 39  6.64      8.92  
6 continentEurope   16.9     0.635       26.7  2.43e-131 15.7      18.2  
7 continentOceania  17.6      1.78       9.92  1.43e- 22  14.2      21.1
```


Tidy regression output with **broom**

```
out_conf %>
  filter(term %nin% "(Intercept)") %>
  mutate(nicelabs = prefix_strip(term, "continent")) %>
  relocate(nicelabs)

# A tibble: 6 × 8
  nicelabs term      estimate std.error statistic   p.value conf.low conf.high
  <chr>    <chr>     <dbl>     <dbl>     <dbl>     <dbl>     <dbl>     <dbl>
1 gdpPercap gdpPercap  4.56e-4  0.0000234    19.4    3.98e- 76  4.10e-4  0.000502
2 log(Pop)  log(pop)   6.39e-1  0.133       4.81    1.64e- 6  3.79e-1  0.900 
3 Americas  continent... 1.31e+1  0.606      21.6    1.85e- 91  1.19e+1  14.3  
4 Asia      continent... 7.78e+0  0.581      13.4    5.52e- 39  6.64e+0  8.92  
5 Europe    continent... 1.69e+1  0.635      26.7    2.43e-131 1.57e+1  18.2  
6 Oceania   continent... 1.76e+1  1.78      9.92    1.43e- 22  1.42e+1  21.1
```

Tidy regression output with **broom**

```
out_conf %>  
  filter(term %nin% "(Intercept)") %>  
  mutate(nicelabs = prefix_strip(term, "continent"))  
ggplot(mapping = aes(x = estimate,  
                      xmin = conf.low,  
                      xmax = conf.high,  
                      y = reorder(nicelabs,  
                                  estimate))) +  
  geom_pointrange() +  
  labs(x = "Estimate",  
       y = NULL,  
       title = "Severely Misspecified")
```



Three ways to tidy

Component level: `tidy()`

Three ways to tidy

Component level: `tidy()`

Observation level: `augment()`

Three ways to tidy

Component level: `tidy()`

Observation level: `augment()`

Model level: `glance()`

Component level

```
> summary(out)

Call:
lm(formula = lifeExp ~ gdpPercap + log(pop) + continent, data = gapminder)
```

Residuals:

Min	1Q	Median	3Q	Max
-47.490	-4.614	0.250	5.293	26.094

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	3.816e+01	2.050e+00	18.618	< 2e-16 ***
gdpPercap	4.557e-04	2.345e-05	19.435	< 2e-16 ***
log(pop)	6.394e-01	1.329e-01	4.810	1.64e-06 ***
continentAmericas	1.308e+01	6.063e-01	21.579	< 2e-16 ***
continentAsia	7.784e+00	5.810e-01	13.398	< 2e-16 ***
continentEurope	1.695e+01	6.350e-01	26.691	< 2e-16 ***
continentOceania	1.764e+01	1.779e+00	9.916	< 2e-16 ***

Signif. codes:	0 '***'	0.001 '**'	0.01 '*'	0.05 '.'
	0.1 ' '	1		

Residual standard error: 8.336 on 1697 degrees of freedom

Multiple R-squared: 0.585, Adjusted R-squared: 0.5835

F-statistic: 398.7 on 6 and 1697 DF, p-value: < 2.2e-16

Observation level

```
> summary(out)

Call:
lm(formula = lifeExp ~ gdpPercap + log(pop) + continent, data = gapminder)

Residuals:
    Min      1Q  Median      3Q     Max 
-47.490 -4.614  0.250  5.293 26.094 

Coefficients:
              Estimate Std. Error t value Pr(>|t|)    
(Intercept) 3.816e+01 2.050e+00 18.618 < 2e-16 ***  
gdpPercap   4.557e-04 2.345e-05 19.435 < 2e-16 ***  
log(pop)    6.394e-01 1.329e-01  4.810 1.64e-06 ***  
continentAmericas 1.308e+01 6.063e-01 21.579 < 2e-16 ***  
continentAsia  7.784e+00 5.810e-01 13.398 < 2e-16 ***  
continentEurope 1.695e+01 6.350e-01 26.691 < 2e-16 ***  
continentOceania 1.764e+01 1.779e+00  9.916 < 2e-16 ***  
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 8.336 on 1697 degrees of freedom
Multiple R-squared:  0.585, Adjusted R-squared:  0.5835 
F-statistic: 398.7 on 6 and 1697 DF,  p-value: < 2.2e-16
```

Observation level

```
augment(out)
```

```
# A tibble: 1,704 × 10
  lifeExp gdpPercap `log(pop)` continent .fitted .resid     .hat .sigma .cooksD
  <dbl>      <dbl>        <dbl> <fct>       <dbl>    <dbl>    <dbl>    <dbl>    <dbl>
1 28.8       779.       15.9 Asia        56.5   -27.7 0.00302  8.31 0.00479
2 30.3       821.       16.0 Asia        56.6   -26.2 0.00299  8.31 0.00426
3 32.0       853.       16.1 Asia        56.7   -24.7 0.00296  8.32 0.00372
4 34.0       836.       16.3 Asia        56.7   -22.7 0.00294  8.32 0.00313
5 36.1       740.       16.4 Asia        56.8   -20.7 0.00294  8.32 0.00259
6 38.4       786.       16.5 Asia        56.9   -18.4 0.00292  8.33 0.00205
7 39.9       978.       16.4 Asia        56.9   -17.0 0.00291  8.33 0.00174
8 40.8       852.       16.4 Asia        56.9   -16.0 0.00292  8.33 0.00155
9 41.7       649.       16.6 Asia        56.9   -15.2 0.00294  8.33 0.00140
10 41.8       635.       16.9 Asia        57.1   -15.3 0.00297  8.33 0.00144
# i 1,694 more rows
# i 1 more variable: .std.resid <dbl>
```

Observation level

For OLS models:

`.fitted` – The fitted values of the model.

`.se.fit` – The standard errors of the fitted values.

`.resid` – The residuals.

`.hat` – The diagonal of the hat matrix.

`.sigma` – An estimate of the residual standard deviation when the corresponding observation is dropped from the model.

`.cooksdist` – Cook's distance, a common regression diagnostic.

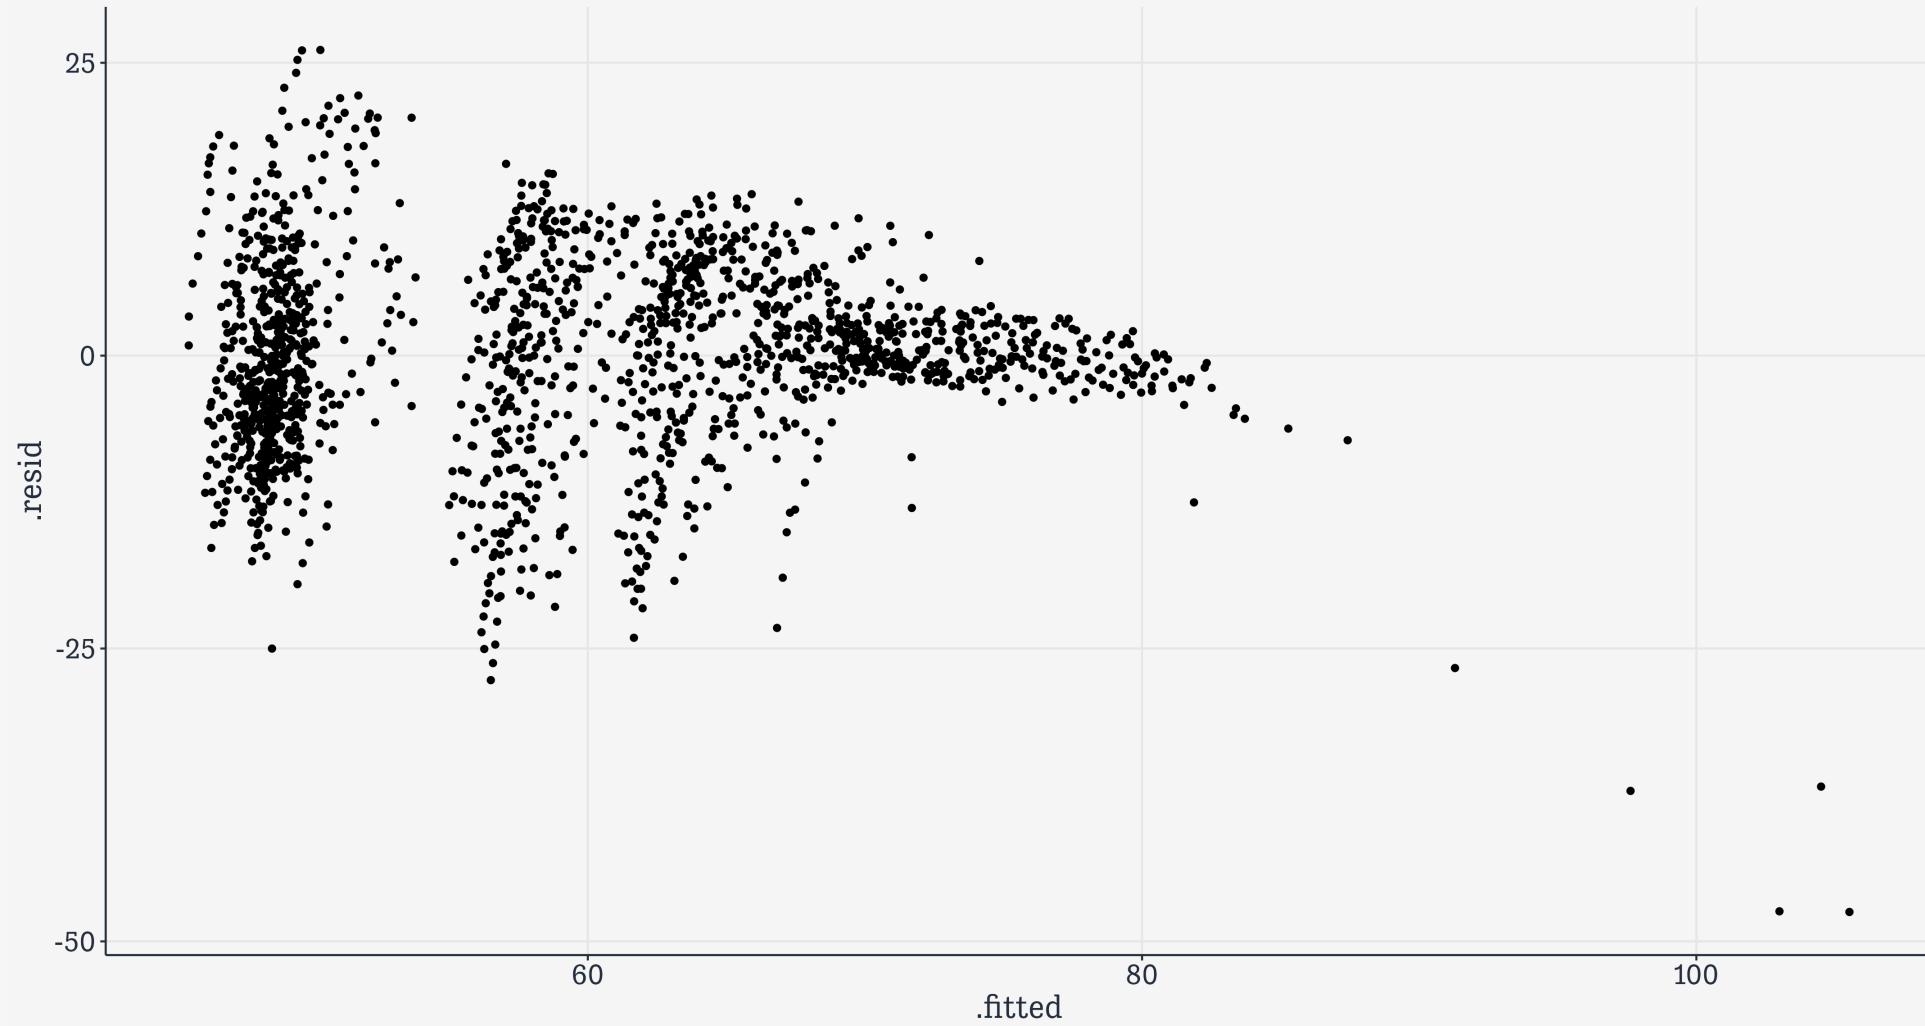
`.std.resid` – The standardized residuals.

Observation level

```
# Adding the data argument puts back any additional columns from the original
# tibble
out_aug ← augment(out, data = gapminder)
head(out_aug)

# A tibble: 6 × 12
  country continent year lifeExp     pop gdpPercap .fitted .resid     .hat .sigma
  <fct>    <fct>   <int>   <dbl>   <int>      <dbl>   <dbl>   <dbl>   <dbl>
1 Afghan... Asia       1952     28.8 8.43e6      779.    56.5  -27.7  0.00302  8.31
2 Afghan... Asia       1957     30.3 9.24e6      821.    56.6  -26.2  0.00299  8.31
3 Afghan... Asia       1962     32.0 1.03e7      853.    56.7  -24.7  0.00296  8.32
4 Afghan... Asia       1967     34.0 1.15e7      836.    56.7  -22.7  0.00294  8.32
5 Afghan... Asia       1972     36.1 1.31e7      740.    56.8  -20.7  0.00294  8.32
6 Afghan... Asia       1977     38.4 1.49e7      786.    56.9  -18.4  0.00292  8.33
# i 2 more variables: .cooksrd <dbl>, .std.resid <dbl>
```

```
## Residuals vs Fitted Values
p ← ggplot(data = out_aug,
            mapping = aes(x = .fitted, y = .resid))
p_out ← p + geom_point()
```



(I told you it was misspecified)

Model level

```
> summary(out)

Call:
lm(formula = lifeExp ~ gdpPercap + log(pop) + continent, data = gapminder)

Residuals:
    Min      1Q  Median      3Q     Max 
-47.490 -4.614  0.250  5.293 26.094 

Coefficients:
            Estimate Std. Error t value Pr(>|t|)    
(Intercept) 3.816e+01 2.050e+00 18.618 < 2e-16 ***  
gdpPercap   4.557e-04 2.345e-05 19.435 < 2e-16 ***  
log(pop)    6.394e-01 1.329e-01  4.810 1.64e-06 ***  
continentAmericas 1.308e+01 6.063e-01 21.579 < 2e-16 ***  
continentAsia    7.784e+00 5.810e-01 13.398 < 2e-16 ***  
continentEurope   1.695e+01 6.350e-01 26.691 < 2e-16 ***  
continentOceania  1.764e+01 1.779e+00  9.916 < 2e-16 ***  
---
Signif. codes:  0 ‘***’ 0.001 ‘**’ 0.01 ‘*’ 0.05 ‘.’ 0.1 ‘ ’ 1
```

```
Residual standard error: 8.336 on 1697 degrees of freedom
```

```
Multiple R-squared:  0.585, Adjusted R-squared:  0.5835
```

```
F-statistic: 398.7 on 6 and 1697 DF,  p-value: < 2.2e-16
```

Model level

```
glance(out)

# A tibble: 1 × 12
  r.squared adj.r.squared sigma statistic   p.value     df logLik     AIC     BIC
      <dbl>         <dbl>    <dbl>     <dbl>     <dbl>     <dbl>    <dbl>    <dbl>
1     0.585        0.584    8.34     399. 1.01e-319     6 -6028. 12072. 12115.
# i 3 more variables: deviance <dbl>, df.residual <int>, nobs <int>
```

The usefulness of `glance()` becomes clearer when dealing with ensembles of models.

Example

A Kaplan-Meier Curve

```
library(survival)
```

```
head(lung)
```

	inst	time	status	age	sex	ph.ecog	ph.karno	pat.karno	meal.cal	wt.loss
1	3	306	2	74	1	1	90	100	1175	NA
2	3	455	2	68	1	0	90	90	1225	15
3	3	1010	1	56	1	0	90	90	NA	15
4	5	210	2	57	1	1	90	60	1150	11
5	1	883	2	60	1	0	100	90	NA	0
6	12	1022	1	74	1	1	50	80	513	0

```
tail(lung)
```

	inst	time	status	age	sex	ph.ecog	ph.karno	pat.karno	meal.cal	wt.loss
223	1	116	2	76	1	1	80	80	NA	0
224	1	188	1	77	1	1	80	60	NA	3
225	13	191	1	39	1	0	90	90	2350	-5
226	32	105	1	75	2	2	60	70	1025	5
227	6	174	1	66	1	1	90	100	1075	1
228	22	177	1	58	2	1	80	90	1060	0

A Kaplan-Meier Curve

First we fit the model:

```
## Hazard model  
out_cph ← coxph(Surv(time, status) ~ age + sex, data = lung)  
  
summary(out_cph)
```

```
Call:  
coxph(formula = Surv(time, status) ~ age + sex, data = lung)  
  
n= 228, number of events= 165  
  
      coef exp(coef)  se(coef)      z Pr(>|z|)  
age   0.017045  1.017191  0.009223  1.848  0.06459 .  
sex  -0.513219  0.598566  0.167458 -3.065  0.00218 **  
---  
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
  
      exp(coef) exp(-coef) lower .95 upper .95  
age     1.0172     0.9831    0.9990    1.0357  
sex     0.5986     1.6707    0.4311    0.8311  
  
Concordance= 0.603  (se = 0.025 )  
Likelihood ratio test= 14.12  on 2 df,  p=9e-04  
Wald test             = 13.47  on 2 df,  p=0.001  
Score (logrank) test = 13.72  on 2 df,  p=0.001
```

A Kaplan-Meier Curve

Then we create the survival curve, which is *nearly* tidy out of the box:

```
## Hazard model
out_surv ← survfit(out_cph)

## See how this is just a print method,
## not a tibble, or even a data frame.
## So it just runs off the end of the slide.
summary(out_surv)
```

Call: survfit(formula = out_cph)

time	n.risk	n.event	survival	std.err	lower	95% CI	upper	95% CI
5	228	1	0.9958	0.00417	0.9877	1.000		
11	227	3	0.9833	0.00831	0.9671	1.000		
12	224	1	0.9791	0.00928	0.9611	0.997		
13	223	2	0.9706	0.01096	0.9494	0.992		
15	221	1	0.9664	0.01170	0.9438	0.990		
26	220	1	0.9622	0.01240	0.9382	0.987		
30	219	1	0.9579	0.01305	0.9327	0.984		
31	218	1	0.9537	0.01368	0.9273	0.981		
53	217	2	0.9452	0.01484	0.9165	0.975		
54	215	1	0.9409	0.01538	0.9112	0.972		
59	214	1	0.9366	0.01590	0.9060	0.968		
60	213	2	0.9280	0.01689	0.8955	0.962		
61	211	1	0.9237	0.01735	0.8903	0.958		
62	210	1	0.9194	0.01780	0.8852	0.955		

[reached getOption("max.print") -- omitted 125 rows]

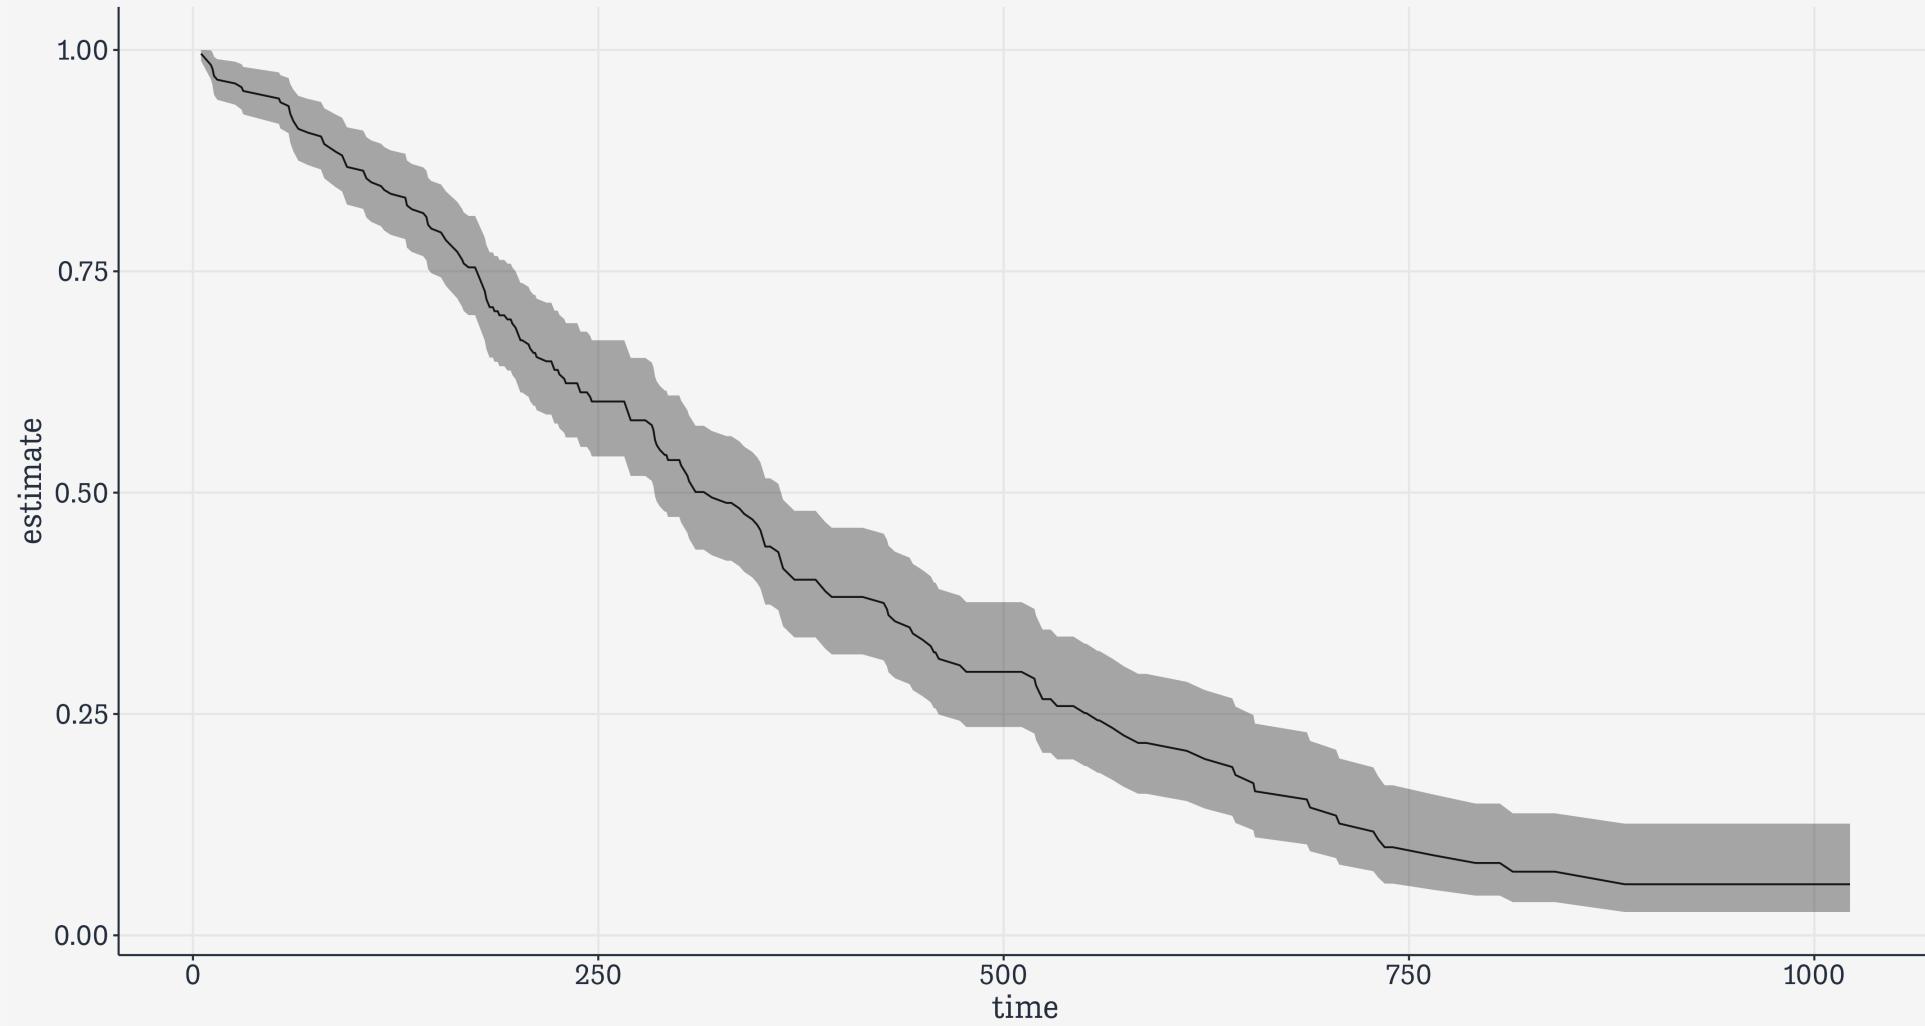
A Kaplan-Meier Curve

Then we tidy it and draw the plot.

```
## Much nicer. (See how the column headers have been regularized, too.)
out_tidy ← tidy(out_surv)
out_tidy

# A tibble: 186 × 8
  time n.risk n.event n.censor estimate std.error conf.high conf.low
  <dbl>   <dbl>   <dbl>   <dbl>     <dbl>    <dbl>      <dbl>    <dbl>
1     5     228     1     0     0.996  0.00419     1     0.988
2    11     227     3     0     0.983  0.00845     1.00    0.967
3    12     224     1     0     0.979  0.00947     0.997    0.961
4    13     223     2     0     0.971  0.0113     0.992    0.949
5    15     221     1     0     0.966  0.0121     0.990    0.944
6    26     220     1     0     0.962  0.0129     0.987    0.938
7    30     219     1     0     0.958  0.0136     0.984    0.933
8    31     218     1     0     0.954  0.0143     0.981    0.927
9    53     217     2     0     0.945  0.0157     0.975    0.917
10   54     215     1     0     0.941  0.0163     0.972    0.911
# i 176 more rows
```

```
p_out ← out_tidy %>
  ggplot(mapping = aes(x = time, y = estimate)) +
  geom_line() +
  geom_ribbon(mapping = aes(ymin = conf.low, ymax = conf.high),
              alpha = 0.4)
```



Kaplan-Meier Plot

Grouped Analysis with broom

Pipelines show
their real power
when used
iteratively

Iteration without tears (or explicit loops)

You might be familiar with code that looks like this:

```
x ← 10  
  
for (i in 1:5) {  
  print(x + i)  
}
```

```
[1] 11  
[1] 12  
[1] 13  
[1] 14  
[1] 15
```

This is one way to do something repeatedly.

Iteration without tears (or explicit loops)

We can also write, e.g.,

```
x ← c(10, 20, 30, 40)

for (i in 1:length(x)) {
  # Add 5 to the ith element of x
  print(x[i] + 5)
}
```

```
[1] 15
[1] 25
[1] 35
[1] 45
```

This way we can refer to each element of `x` in turn, and do the same thing to it.

Iteration without tears (or explicit loops)

The more complicated the thing we want to do, the more likely we are to use functions to help us out.

```
x ← 10  
  
for (i in 1:5) {  
  print(sqrt(x + i))  
}
```

```
[1] 3.316625  
[1] 3.464102  
[1] 3.605551  
[1] 3.741657  
[1] 3.872983
```

Isn't this like ... Vectorized arithmetic?

The simplest cases are not that different from the vectorized arithmetic we saw before.

```
a ← c(1:10)  
  
b ← 1  
  
# You know what R will do here  
a + b
```

```
[1]  2  3  4  5  6  7  8  9 10 11
```

Isn't this like ... Vectorized arithmetic?

The simplest cases are not that different from the vectorized arithmetic we saw before.

```
a ← c(1:10)  
  
b ← 1  
  
# You know what R will do here  
a + b
```

[1] 2 3 4 5 6 7 8 9 10 11

R's vectorized rules add **b** to every element of **a**. In a sense, the **+** operation can be thought of as a function that takes each element of **a** and does something with it. In this case “add **b**”.

Repeatedly applying a function

We can make this explicit by writing a function:

```
a ← c(1:10)

add_b ← function(x) {
  b ← 1
  x + b # for any x
}
```

Now:

```
add_b(x = a)
[1]  2  3  4  5  6  7  8  9 10 11
```

In effect we take the vector **a** and feed it to the **add_b()** function one element at a time.

Repeatedly applying a function

Again, R's vectorized approach means it automatically applies `add_b()` to every element of the `x` we give it.

```
add_b(x = 10)
```

```
[1] 11
```

```
add_b(x = c(1, 99, 1000))
```

```
[1] 2 100 1001
```

Iterating in a pipeline

Some operations can't directly be vectorized in this way, most often because the function we want to apply only knows what to do if it is handed, say, a vector. It doesn't understand what to do if it's handed a list of vectors or a tibble of them, etc. This is when we might find ourselves manually iterating—writing out every single step explicitly.

```
library(gapminder)
gapminder %>%
  summarize(country_n = n_distinct(country),
           continent_n = n_distinct(continent),
           year_n = n_distinct(year),
           lifeExp_n = n_distinct(lifeExp),
           population_n = n_distinct(population))

# A tibble: 1 × 5
  country_n continent_n year_n lifeExp_n population_n
  <int>      <int>   <int>     <int>        <int>
1       142          5     12      1626        4060
```

That's tedious to write! Computers are supposed to allow us to avoid that sort of thing.

Iterating in a pipeline

So how would we iterate this? What we want is to apply the `n_distinct()` function to each column of `gapminder`. But we can't easily write a loop inside a pipeline. We want a way to iterate that lets us repeatedly apply a function without explicitly writing a loop.

```
library(gapminder)
gapminder %>
  summarize(n_distinct(country),
            n_distinct(continent),
            n_distinct(year),
            n_distinct(lifeExp),
            n_distinct(population))

# A tibble: 1 × 5
`n_distinct(country)` `n_distinct(continent)` `n_distinct(year)` 
<int>                  <int>                  <int>
1                      142                      5                  12
# i 2 more variables: `n_distinct(lifeExp)` <int>,
#   `n_distinct(population)` <int>
```

Using `n_distinct()` in this context is an idea I got from Rebecca Barter's discussion of `purrr`.

Iterating in a pipeline

In real life, you'd use **across()**, like this:

```
gapminder >
  summarize(across(everything(), n_distinct))

# A tibble: 1 × 6
  country continent year lifeExp   pop gdpPercap
  <int>     <int> <int>    <int> <int>      <int>
1     142         5    12     1626  1704       1704
```

Iterating in a pipeline

But you could also say “Feed each column of `gapminder` in turn to the `n_distinct()` function”. This is what `map()` is for.

```
map(gapminder, n_distinct)
```

```
$country  
[1] 142
```

```
$continent  
[1] 5
```

```
$year  
[1] 12
```

```
$lifeExp  
[1] 1626
```

```
$pop  
[1] 1704
```

```
$gdpPercap  
[1] 1704
```

Read it as “Feed each column of `gapminder` to the `n_distinct()` function.
(This is pretty much what `across()` is doing more nicely.)

Iterating in a pipeline

Or, in pipeline form:

```
gapminder %>  
  map(n_distinct)
```

```
$country  
[1] 142
```

```
$continent  
[1] 5
```

```
$year  
[1] 12
```

```
$lifeExp  
[1] 1626
```

```
$pop  
[1] 1704
```

```
$gdpPercap  
[1] 1704
```

You can see we are getting a *list* back.

Iterating in a pipeline

Or, in pipeline form:

```
result ← gapminder ▷  
  map(n_distinct)  
  
class(result)  
  
[1] "list"
```

```
result$continent  
  
[1] 5
```

```
result[[2]]  
  
[1] 5
```

Iterating in a pipeline

But we know `n_distinct()` should always return an integer. So we use `map_int()` instead of the generic `map()`.

```
gapminder >
  map_int(n_distinct)
```

country	continent	year	lifeExp	pop	gdpPercap
142	5	12	1626	1704	1704

The thing about the `map()` family is that it can deal with all kinds of input types and output types.

So what's the use of
all that stuff?

Grouped analysis and **list columns**

Let's say I want to fit a model to data for all countries in Europe in 1977.

```
eu77 ← gapminder ▷  
  filter(continent = "Europe", year = 1977)  
  
fit ← lm(lifeExp ~ log(gdpPercap), data = eu77)  
  
summary(fit)
```

```
Call:  
lm(formula = lifeExp ~ log(gdpPercap), data = eu77)  
  
Residuals:  
    Min      1Q  Median      3Q     Max  
-7.4956 -1.0306  0.0935  1.1755  3.7125  
  
Coefficients:  
            Estimate Std. Error t value Pr(>|t|)  
(Intercept)  29.489     7.161   4.118 0.000306 ***  
log(gdpPercap) 4.488     0.756   5.936 2.17e-06 ***  
---  
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
  
Residual standard error: 2.114 on 28 degrees of freedom  
Multiple R-squared:  0.5572,    Adjusted R-squared:  0.5414  
F-statistic: 35.24 on 1 and 28 DF,  p-value: 2.173e-06
```

Grouped analysis and **list columns**

What if I want to do that for all Continent-Year combinations?

```
out_le ← gapminder ▷ group_by(continent, year) ▷  
  nest()
```

```
out_le
```

```
# A tibble: 60 × 3  
# Groups:   continent, year [60]  
  continent   year data  
  <fct>     <int> <list>  
1 Asia        1952 <tibble [33 × 4]>  
2 Asia        1957 <tibble [33 × 4]>  
3 Asia        1962 <tibble [33 × 4]>  
4 Asia        1967 <tibble [33 × 4]>  
5 Asia        1972 <tibble [33 × 4]>  
6 Asia        1977 <tibble [33 × 4]>  
7 Asia        1982 <tibble [33 × 4]>  
8 Asia        1987 <tibble [33 × 4]>  
9 Asia        1992 <tibble [33 × 4]>  
10 Asia       1997 <tibble [33 × 4]>  
# i 50 more rows
```

Think of nesting as a kind of “super-grouping”. Look in the object inspector.

Grouped analysis and **list columns**

Europe '77 is still in there.

```
out_le >
  filter(continent == "Europe" & year == 1977) >
  unnest(cols = c(data))

# A tibble: 30 × 6
# Groups:   continent, year [1]
  continent  year country      lifeExp     pop gdpPercap
  <fct>     <int> <fct>       <dbl>     <int>     <dbl>
1 Europe     1977 Albania      68.9    2509048    3533.
2 Europe     1977 Austria      72.2    7568430    19749.
3 Europe     1977 Belgium      72.8    9821800    19118.
4 Europe     1977 Bosnia and Herzegovina 69.9    4086000    3528.
5 Europe     1977 Bulgaria      70.8    8797022    7612.
6 Europe     1977 Croatia      70.6    4318673    11305.
7 Europe     1977 Czech Republic 70.7    10161915   14800.
8 Europe     1977 Denmark      74.7    5088419    20423.
9 Europe     1977 Finland      72.5    4738902    15605.
10 Europe    1977 France       73.8    53165019   18293.
# i 20 more rows
```

Grouped analysis and **list columns**

Here we write a tiny, very specific function and **map()** it to every row in the **data** column.

```
fit_ols ← function(df) {  
  lm(lifeExp ~ log(gdpPercap), data = df)  
}  
  
out_le ← gapminder ▷  
  group_by(continent, year) ▷  
  nest() ▷  
  mutate(model = map(data, fit_ols))
```

Grouped analysis and **list columns**

Now we have a new column. Each row of the `model` column contains a full regression for that continent-year.

```
out_le
```

```
# A tibble: 60 × 4
# Groups:   continent, year [60]
  continent  year data          model
  <fct>     <int> <list>        <list>
1 Asia       1952 <tibble [33 × 4]> <lm>
2 Asia       1957 <tibble [33 × 4]> <lm>
3 Asia       1962 <tibble [33 × 4]> <lm>
4 Asia       1967 <tibble [33 × 4]> <lm>
5 Asia       1972 <tibble [33 × 4]> <lm>
6 Asia       1977 <tibble [33 × 4]> <lm>
7 Asia       1982 <tibble [33 × 4]> <lm>
8 Asia       1987 <tibble [33 × 4]> <lm>
9 Asia       1992 <tibble [33 × 4]> <lm>
10 Asia      1997 <tibble [33 × 4]> <lm>
# i 50 more rows
```

Grouped analysis and **list columns**

We can tidy the nested models, too.

```
fit_ols ← function(df) {  
  lm(lifeExp ~ log(gdpPerCap), data = df)  
}  
  
out_tidy ← gapminder ▷  
  group_by(continent, year) ▷  
  nest() ▷  
  mutate(model = map(data, fit_ols),  
         tidied = map(model, tidy))  
  
out_tidy  
  
# A tibble: 60 × 5  
# Groups:   continent, year [60]  
  continent  year data          model  tidied  
  <fct>     <int> <list>        <list> <list>  
1 Asia       1952 <tibble [33 × 4]> <lm>    <tibble [2 × 5]>  
2 Asia       1957 <tibble [33 × 4]> <lm>    <tibble [2 × 5]>  
3 Asia       1962 <tibble [33 × 4]> <lm>    <tibble [2 × 5]>  
4 Asia       1967 <tibble [33 × 4]> <lm>    <tibble [2 × 5]>  
5 Asia       1972 <tibble [33 × 4]> <lm>    <tibble [2 × 5]>  
6 Asia       1977 <tibble [33 × 4]> <lm>    <tibble [2 × 5]>  
7 Asia       1982 <tibble [33 × 4]> <lm>    <tibble [2 × 5]>  
8 Asia       1987 <tibble [33 × 4]> <lm>    <tibble [2 × 5]>  
9 Asia       1992 <tibble [33 × 4]> <lm>    <tibble [2 × 5]>  
10 Asia      1997 <tibble [33 × 4]> <lm>    <tibble [2 × 5]>  
# i 50 more rows
```

Grouped analysis and **list columns**

We can get the tidied results out into the main table if we like.

```
out_tidy ← out_tidy %>
  unnest(cols = c(tidied)) %>
  filter(term %nin% "(Intercept)" &
        continent %nin% "Oceania")

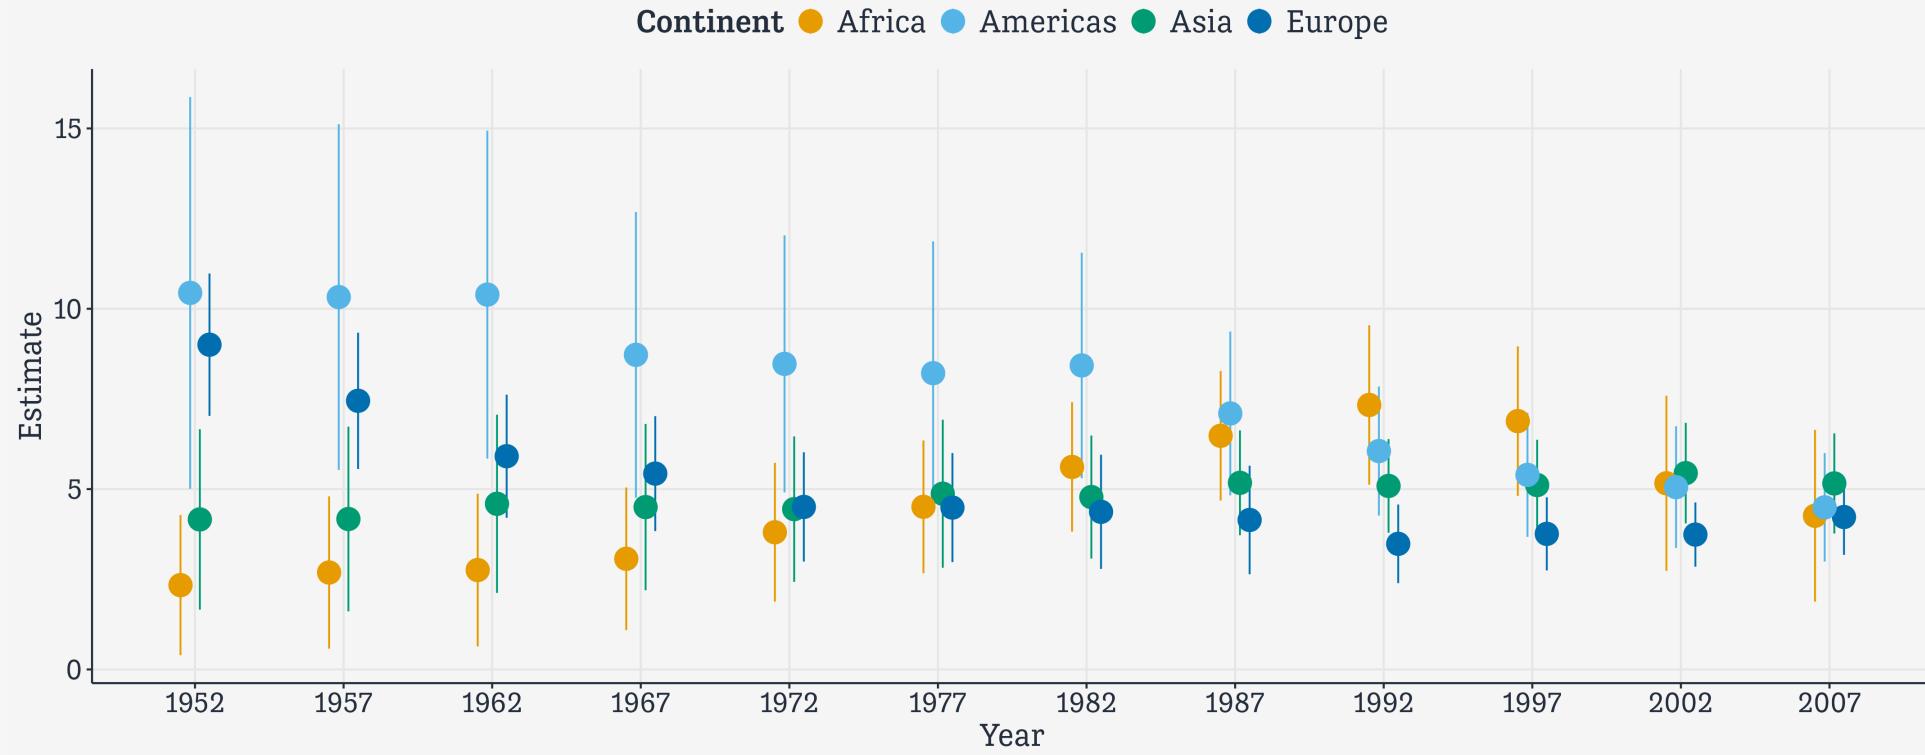
out_tidy
```

A tibble: 48 × 9
Groups: continent, year [48]
continent year data model term estimate std.error statistic p.value
<fct> <int> <tibble> <lm> <chr> <dbl> <dbl> <dbl> <dbl>
1 Asia 1952 <tibble> <lm> log(gdp... 4.16 1.25 3.33 2.28e-3
2 Asia 1957 <tibble> <lm> log(gdp... 4.17 1.28 3.26 2.71e-3
3 Asia 1962 <tibble> <lm> log(gdp... 4.59 1.24 3.72 7.94e-4
4 Asia 1967 <tibble> <lm> log(gdp... 4.50 1.15 3.90 4.77e-4
5 Asia 1972 <tibble> <lm> log(gdp... 4.44 1.01 4.41 1.16e-4
6 Asia 1977 <tibble> <lm> log(gdp... 4.87 1.03 4.75 4.42e-5
7 Asia 1982 <tibble> <lm> log(gdp... 4.78 0.852 5.61 3.77e-6
8 Asia 1987 <tibble> <lm> log(gdp... 5.17 0.727 7.12 5.31e-8
9 Asia 1992 <tibble> <lm> log(gdp... 5.09 0.649 7.84 7.60e-9
10 Asia 1997 <tibble> <lm> log(gdp... 5.11 0.628 8.15 3.35e-9
i 38 more rows

Plot what we have

```
p ← ggplot(data = out_tidy,
             mapping = aes(x = year, y = estimate,
                           ymin = estimate - 2*std.error,
                           ymax = estimate + 2*std.error,
                           group = continent,
                           color = continent))

p_out ← p +
  geom_pointrange(size = rel(1.25),
                  position = position_dodge(width = rel(1.3))) +
  scale_x_continuous(breaks = unique(gapminder$year)) +
  labs(x = "Year",
       y = "Estimate",
       color = "Continent")
```



Repeated Estimates of log GDP on Life Expectancy by Continent

And there's more ...

Let's go back to this stage:

```
# New model
fit_ols2 ← function(df) {
  lm(lifeExp ~ log(gdpPercap) + log(pop), data = df)
}

out_tidy ← gapminder ▷
  group_by(continent, year) ▷
  nest() ▷
  mutate(model = map(data, fit_ols2),
         tidied = map(model, tidy))

out_tidy

# A tibble: 60 × 5
# Groups:   continent, year [60]
  continent  year data          model  tidied
  <fct>     <int> <list>        <list> <list>
1 Asia       1952 <tibble [33 × 4]> <lm>    <tibble [3 × 5]>
2 Asia       1957 <tibble [33 × 4]> <lm>    <tibble [3 × 5]>
3 Asia       1962 <tibble [33 × 4]> <lm>    <tibble [3 × 5]>
4 Asia       1967 <tibble [33 × 4]> <lm>    <tibble [3 × 5]>
5 Asia       1972 <tibble [33 × 4]> <lm>    <tibble [3 × 5]>
6 Asia       1977 <tibble [33 × 4]> <lm>    <tibble [3 × 5]>
7 Asia       1982 <tibble [33 × 4]> <lm>    <tibble [3 × 5]>
8 Asia       1987 <tibble [33 × 4]> <lm>    <tibble [3 × 5]>
9 Asia       1992 <tibble [33 × 4]> <lm>    <tibble [3 × 5]>
10 Asia      1997 <tibble [33 × 4]> <lm>   <tibble [3 × 5]>

# i 50 more rows
```

A function to draw a coef plot

```
# Plot the output from our model
mod_plot ← function(data,
                      title){
  data %>
    filter(term %nin% "(Intercept)") %>
    ggplot(mapping = aes(x = estimate,
                          xmin = estimate - std.error,
                          xmax = estimate + std.error,
                          y = reorder(term, estimate))) +
    geom_pointrange() +
    labs(title = title,
         y = NULL)
}
```


Add it using `map2()` or `pmap()`

When we have two arguments to feed a function we can use `map2()`. The general case is `pmap()`, for passing any number of arguments in a list.


```
out_tidy ← gapminder ▷ group_by(continent, year) ▷ nest() ▷  
  mutate(title = paste(continent, year),  
        model = map(data, fit_ols2),  
        tidied = map(model, tidy),  
        ggout = pmap(list(tidied, title),  
                    mod_plot))  
  
out_tidy  
  
# A tibble: 60 × 7  
# Groups:   continent, year [60]  
  continent  year data           title    model  tidied      ggout  
  <fct>     <int> <list>         <chr>    <list> <list>       <list>  
1 Asia       1952 <tibble [33 × 4]> Asia 1952 <lm>  <tibble [3 × 5]> <gg>  
2 Asia       1957 <tibble [33 × 4]> Asia 1957 <lm>  <tibble [3 × 5]> <gg>  
3 Asia       1962 <tibble [33 × 4]> Asia 1962 <lm>  <tibble [3 × 5]> <gg>  
4 Asia       1967 <tibble [33 × 4]> Asia 1967 <lm>  <tibble [3 × 5]> <gg>  
5 Asia       1972 <tibble [33 × 4]> Asia 1972 <lm>  <tibble [3 × 5]> <gg>  
6 Asia       1977 <tibble [33 × 4]> Asia 1977 <lm>  <tibble [3 × 5]> <gg>  
7 Asia       1982 <tibble [33 × 4]> Asia 1982 <lm>  <tibble [3 × 5]> <gg>  
8 Asia       1987 <tibble [33 × 4]> Asia 1987 <lm>  <tibble [3 × 5]> <gg>  
9 Asia       1992 <tibble [33 × 4]> Asia 1992 <lm>  <tibble [3 × 5]> <gg>  
10 Asia      1997 <tibble [33 × 4]> Asia 1997 <lm>  <tibble [3 × 5]> <gg>  
# i 50 more rows
```

A plot!

```
out_tidy$ggout[[8]]
```



```
out_tidy$ggout[[18]]
```



We don't just put them in there for fun

We can e.g. `walk` the plots out to disk

`walk()` is `map()` for when you just want a “side-effect” such as printed output.
There is also `walk2()` and `pwalk()`

```
pwalk(  
  list(  
    filename = paste0(out_tidy$title, ".png"),  
    plot = out_tidy$ggout,  
    path = here("figures"),  
    height = 3, width = 4,  
    dpi = 300  
  ),  
  ggsave  
)
```

Peek in the `figures/` folder

```
fs::dir_ls(here("figures")) %>  
  basename()  
  
[1] "Africa 1952.png"    "Africa 1957.png"    "Africa 1962.png"  
[4] "Africa 1967.png"    "Africa 1972.png"    "Africa 1977.png"  
[7] "Africa 1982.png"    "Africa 1987.png"    "Africa 1992.png"  
[10] "Africa 1997.png"   "Africa 2002.png"   "Africa 2007.png"  
[13] "Americas 1952.png" "Americas 1957.png" "Americas 1962.png"  
[16] "Americas 1967.png" "Americas 1972.png" "Americas 1977.png"  
[19] "Americas 1982.png" "Americas 1987.png" "Americas 1992.png"  
[22] "Americas 1997.png" "Americas 2002.png" "Americas 2007.png"  
[25] "Asia 1952.png"     "Asia 1957.png"     "Asia 1962.png"  
[28] "Asia 1967.png"     "Asia 1972.png"     "Asia 1977.png"  
[31] "Asia 1982.png"     "Asia 1987.png"     "Asia 1992.png"  
[34] "Asia 1997.png"     "Asia 2002.png"     "Asia 2007.png"  
[37] "Europe 1952.png"   "Europe 1957.png"   "Europe 1962.png"  
[40] "Europe 1967.png"   "Europe 1972.png"   "Europe 1977.png"  
[43] "Europe 1982.png"   "Europe 1987.png"   "Europe 1992.png"  
[46] "Europe 1997.png"   "Europe 2002.png"   "Europe 2007.png"  
[49] "Oceania 1952.png"  "Oceania 1957.png"  "Oceania 1962.png"  
[52] "Oceania 1967.png"  "Oceania 1972.png"  "Oceania 1977.png"  
[55] "Oceania 1982.png"  "Oceania 1987.png"  "Oceania 1992.png"
```

Get model-based
graphics right

Present findings in substantive terms

Show degrees of
confidence or
uncertainty

Show the data
when you can

But all this applies
to data in *any*
format. Graphs are
not special here!

**Plot Marginal Effects with the
marginaleffects package**

An example from the GSS

```
gss_sm
```

```
# A tibble: 2,867 × 32
  year   id ballot      age child� sibs degree race   sex   region income16
  <dbl> <dbl> <labelled> <dbl> <dbl> <labe> <fct> <fct> <fct> <fct> <fct>
1 2016    1 1           47     3 2  Bach... White Male  New E... $170000...
2 2016    2 2           61     0 3  High ... White Male  New E... $50000 ...
3 2016    3 3           72     2 3  Bach... White Male  New E... $75000 ...
4 2016    4 1           43     4 3  High ... White Fema... New E... $170000...
5 2016    5 3           55     2 2  Gradu... White Fema... New E... $170000...
6 2016    6 2           53     2 2  Junio... White Fema... New E... $60000 ...
7 2016    7 1           50     2 2  High ... White Male  New E... $170000...
8 2016    8 3           23     3 6  High ... Other Fema... Middl... $30000 ...
9 2016    9 1           45     3 5  High ... Black Male  Middl... $60000 ...
10 2016   10 3          71     4 1  Junio... White Male  Middl... $60000 ...
# i 2,857 more rows
# i 21 more variables: relig <fct>, marital <fct>, padege <fct>, madeg <fct>,
# partyid <fct>, polviews <fct>, happy <fct>, partners <fct>, grass <fct>,
# zodiac <fct>, pres12 <labelled>, wtssall <dbl>, income_rc <fct>,
# agegrp <fct>, ageq <fct>, siblings <fct>, kids <fct>, religion <fct>,
# bigregion <fct>, partners_rc <fct>, obama <dbl>
```

Set up our model

```
gss_sm$polviews_m ← relevel(gss_sm$polviews,
                             ref = "Moderate")

out_bo ← glm(obama ~ polviews_m + sex*race,
             family = "binomial",
             data = gss_sm)

tidy(out_bo)

# A tibble: 12 × 5
  term                  estimate std.error statistic p.value
  <chr>                 <dbl>     <dbl>      <dbl>    <dbl>
1 (Intercept)           0.296     0.134      2.21    2.70e- 2
2 polviews_mExtremely Liberal  2.37      0.525      4.52    6.20e- 6
3 polviews_mLiberal     2.60      0.357      7.29    3.10e-13
4 polviews_mSlightly Liberal  1.29      0.248      5.21    1.94e- 7
5 polviews_mSlightly Conservative -1.36     0.181     -7.48   7.68e-14
6 polviews_mConservative   -2.35     0.200     -11.7    1.07e-31
7 polviews_mExtremely Conservative -2.73     0.387     -7.04   1.87e-12
8 sexFemale              0.255     0.145      1.75    7.96e- 2
9 raceBlack               3.85      0.501      7.68    1.61e-14
10 raceOther              -0.00214    0.436     -0.00492 9.96e- 1
11 sexFemale:raceBlack     -0.198     0.660     -0.299   7.65e- 1
12 sexFemale:raceOther      1.57      0.588      2.68    7.37e- 3
```

Calculate the Average Marginal Effects

```
library(marginaleffects)

bo_mfx ← avg_slopes(out_bo)

## This gives us the marginal effects at the unit level
as_tibble(bo_mfx)

# A tibble: 9 × 9
  term      contrast   estimate std.error statistic   p.value s.value conf.low
  <chr>    <chr>        <dbl>     <dbl>      <dbl>      <dbl>    <dbl>    <dbl>
1 polviews_m Conservati... -0.412     0.0283     -14.5  6.82e- 48  157.    -0.467
2 polviews_m Extremely ... -0.454     0.0420     -10.8  3.55e- 27  87.9    -0.536
3 polviews_m Extremely ...  0.268     0.0295      9.10  9.07e- 20  63.3    0.210
4 polviews_m Liberal - ...  0.277     0.0229     12.1  1.46e- 33  109.    0.232
5 polviews_m Slightly C... -0.266     0.0330     -8.06 7.65e- 16  50.2    -0.330
6 polviews_m Slightly L...  0.193     0.0303      6.39  1.66e- 10  32.5    0.134
7 race       Black - Wh...  0.403     0.0173     23.4  1.18e-120  398.    0.369
8 race       Other - Wh...  0.125     0.0386      3.23  1.24e-  3  9.66    0.0490
9 sex        Female - M...  0.0443    0.0177      2.51  1.22e-  2  6.36    0.00967
# i 1 more variable: conf.high <dbl>
```


Alternatively, do it with **broom**

```
tidy(bo_mfx)
```

```
# A tibble: 9 × 9
  term      contrast estimate std.error statistic p.value s.value conf.low
  <chr>     <chr>     <dbl>     <dbl>     <dbl>     <dbl>     <dbl>     <dbl>
1 polviews_m Conservati... -0.412    0.0283   -14.5    6.82e- 48  157.    -0.467
2 polviews_m Extremely ... -0.454    0.0420   -10.8    3.55e- 27  87.9    -0.536
3 polviews_m Extremely ...  0.268    0.0295    9.10    9.07e- 20  63.3    0.210
4 polviews_m Liberal - ...  0.277    0.0229   12.1    1.46e- 33  109.    0.232
5 polviews_m Slightly C... -0.266    0.0330   -8.06    7.65e- 16  50.2    -0.330
6 polviews_m Slightly L...  0.193    0.0303    6.39    1.66e- 10  32.5    0.134
7 race       Black - Wh...  0.403    0.0173   23.4    1.18e-120 398.    0.369
8 race       Other - Wh...  0.125    0.0386    3.23    1.24e-  3  9.66    0.0490
9 sex        Female - M...  0.0443   0.0177    2.51    1.22e-  2  6.36    0.00967
# i 1 more variable: conf.high <dbl>
```

Which gets us back to familiar territory

```
tidy(bo_mfx) %>%  
  ggplot(mapping = aes(x = estimate,  
                        xmin = conf.low,  
                        xmax = conf.high,  
                        y = reorder(contrast,  
                                    estimate))) +  
  geom_vline(xintercept = 0, color = "gray70",  
             size = rel(1.2)) +  
  geom_pointrange() +  
  labs(x = "Average Marginal Effect",  
       y = NULL)
```



marginaleffects can do a lot more

marginaleffects

0.4.1.9000

Adjusted predictions

Marginal effects

Contrasts

Marginal means

Functions

Changelog

The marginaleffects package for R

marginaleffects is an R package to compute and plot adjusted predictions, marginal effects, contrasts, and marginal means for a *wide* variety of models.



Links

[View on CRAN](#)

[Browse source code](#)

[Report a bug](#)

License

[Full license](#)

GPL (>= 3)

Citation

[Citing marginaleffects](#)

Developers

Vincent Arel-Bundock

Author, maintainer, copyright holder

Dev status

codecov 90%

R-CMD-check passing

CRAN 0.4.1

dependencies 3/4

Table of contents

- [What?](#)
- [Why?](#)
- [Getting started](#)
- Vignettes:
 - [Adjusted predictions](#)
 - [Marginal effects](#)
 - [Contrasts](#)
 - [Marginal means](#)
- [61 supported models](#)
- [Performance tips](#)
- Case studies:
 - [Bayesian analyses with brms](#)
 - [Mixed effects models](#)
 - [Generalized Additive Models](#)
 - [Multinomial Logit and Discrete Choice Models](#)
 - [Tables and plots](#)
 - [Robust standard errors and more](#)
 - [Transformations and Custom Contrasts: Adjusted Risk Ratio Example](#)
- [Alternative software](#)
- [Technical notes](#)

marginaleffects can do a lot more

It includes a range of plotting methods, to produce graphics directly.

These are built on `ggplot`. Similarly integration with `broom` means that you can use the package-specific plotting functions take the tidy output and adapt it to your own needs.

Also check out `modelsummary`, by the same author, for quick and flexible summaries of models and datasets. Again, this sort of package is very convenient to use directly. But with just a little facility with R and tidyverse-style idioms and patterns, you'll get even more out of it. You'll better understand how to adapt it and why its functions work as they do.

Complex Surveys with the survey and srvyr packages

Working with complex surveys

As always, our question is “What’s the smoothest way for me to get a **tidy table of results** I need to hand off to **ggplot**?“

For complex surveys, we use **survey**, the standard package for survey analysis in R, and **srvyr**, a helper package designed to integrate what **survey** can do with the Tidyverse framework.

```
## Load the packages
library(survey)
library(srvyr)
```

Example: The GSS again

This time, a small piece of the full GSS from the early 1970s to 2018.

```
gss_lon
```

```
# A tibble: 62,466 × 25
  year     id ballot age degree race   sex siblings kids bigregion income16
  <dbl> <dbl> <labe> <lab> <fct> <fct> <fct> <fct> <fct> <fct>
1 1972      1 NA    23  Bache... White Fema... 3      0 Midwest <NA>
2 1972      2 NA    70  Lt Hi... White Male  4     4+ Midwest <NA>
3 1972      3 NA    48  High ... White Fema... 5     4+ Midwest <NA>
4 1972      4 NA    27  Bache... White Fema... 5      0 Midwest <NA>
5 1972      5 NA    61  High ... White Fema... 2      2 Midwest <NA>
6 1972      6 NA    26  High ... White Male  1      0 Midwest <NA>
7 1972      7 NA    28  High ... White Male  6+     2 Midwest <NA>
8 1972      8 NA    27  Bache... White Male  1      0 Midwest <NA>
9 1972      9 NA    21  High ... Black Fema... 2      2 South   <NA>
10 1972     10 NA   30  High ... Black Fema... 6+     4+ South   <NA>
# i 62,456 more rows
# i 14 more variables: religion <fct>, marital <fct>, padeg <fct>, madeg <fct>,
# partyid <fct>, polviews <fct>, happy <fct>, partners_rc <fct>, grass <fct>,
# zodiac <fct>, pres12 <labelled>, wtssall <dbl>, vpsu <dbl>, vstrat <dbl>
```

Add the weighting information

```
# These details are dependent on the kind of survey you're working with
options(survey.lonely.psu = "adjust")
options(na.action="na.pass")

gss_svy ← gss_lon ▷
  filter(year > 1974) ▷
  mutate(stratvar = interaction(year, vstrat)) ▷
  as_survey_design(ids = vpsu,
                   strata = stratvar,
                   weights = wtssall,
                   nest = TRUE)

gss_svy # Now it's no longer simply a tibble
```

```
Stratified 1 - level Cluster Sampling design (with replacement)
With (4399) clusters.
Called via srvyr
Sampling variables:
- ids: vpsu
- strata: stratvar
- weights: wtssall
Data variables:
- year (dbl), id (dbl), ballot (labelled), age (labelled), degree (fct), race
(fct), sex (fct), siblings (fct), kids (fct), bigregion (fct), income16
(fct), religion (fct), marital (fct), padeg (fct), madeg (fct), partyid
(fct), polviews (fct), happy (fct), partners_rc (fct), grass (fct), zodiac
(fct), pres12 (labelled), wtssall (dbl), vpsu (dbl), vstrat (dbl), stratvar
(fct)
```


Trends in the happy measure

```
out_hap ← gss_svy ▷  
  group_by(year, happy) ▷  
  summarize(prop = survey_mean(na.rm = TRUE, vartype = "ci"))  
  
out_hap  
  
# A tibble: 111 × 5  
# Groups:   year [28]  
  year    happy      prop  prop_low  prop_upp  
  <dbl> <fct>     <dbl>    <dbl>     <dbl>  
1 1975 Very Happy  0.333    0.304    0.361  
2 1975 Pretty Happy 0.534    0.504    0.565  
3 1975 Not Too Happy 0.129    0.108    0.150  
4 1975 <NA>        0.00375  0.000362  0.00714  
5 1976 Very Happy  0.348    0.321    0.376  
6 1976 Pretty Happy 0.529    0.504    0.555  
7 1976 Not Too Happy 0.122    0.104    0.141  
8 1977 Very Happy  0.357    0.330    0.384  
9 1977 Pretty Happy 0.532    0.504    0.560  
10 1977 Not Too Happy 0.110   0.0918   0.128  
# i 101 more rows
```


Once again, it's now a tidy tibble, and we know what to do with those.

Trends in the happy measure

```
out_hap >
  filter(happy == "Not Too Happy") >
  ggplot(mapping = aes(x = year,
                        y = prop,
                        ymin = prop_low,
                        ymax = prop_upp)) +
  geom_line(linewidth = 1.2) +
  geom_ribbon(alpha = 0.3) +
  scale_x_continuous(breaks =
    seq(1978, 2018, 4)) +
  scale_y_continuous(labels =
    label_percent(accuracy = 1)) +
  labs(x = "Year",
       y = "Percent",
       title = "Trends in Unhappiness",
       subtitle = "1975-2018",
       caption = "Data: GSS.")
```



With a better aspect ratio



A more complex example

```
gss_svy %>
  filter(year %in% seq(1976, 2016, by = 4)) %>
  group_by(year, race, degree) %>
  summarize(prop = survey_mean(na.rm = TRUE))

# A tibble: 162 × 5
# Groups:   year, race [30]
  year race  degree      prop    prop_se
  <dbl> <fct> <fct>     <dbl>    <dbl>
1 1976 White Lt High School 0.327    0.0160
2 1976 White High School   0.517    0.0161
3 1976 White Junior College 0.0128   0.00298
4 1976 White Bachelor     0.101    0.00955
5 1976 White Graduate     0.0392   0.00642
6 1976 White <NA>         0.00285  0.00151
7 1976 Black Lt High School 0.558    0.0603
8 1976 Black High School   0.335    0.0476
9 1976 Black Junior College 0.0423   0.0192
10 1976 Black Bachelor    0.0577   0.0238
# i 152 more rows
```

Let's put that in an object

```
out_yrd ← gss_svy ▷  
  filter(year %in% seq(1976, 2016, by = 4)) ▷  
  group_by(year, race, degree) ▷  
  summarize(prop = survey_mean(na.rm = TRUE))
```

Check the sums

```
out_yrd %>  
  group_by(year, race) %>  
  summarize(tot = sum(prop))
```


```
# A tibble: 30 × 3  
# Groups:   year [10]  
  year race   tot  
  <dbl> <fct> <dbl>  
1 1976 White  1.00  
2 1976 Black  1.00  
3 1976 Other  1  
4 1980 White  1.00  
5 1980 Black  1  
6 1980 Other  1  
7 1984 White  1.00  
8 1984 Black  1.00  
9 1984 Other  1  
10 1988 White  1.00  
# i 20 more rows
```

Set up the plot

```
p ← out_yrd ▷  
  drop_na() ▷  
  filter(race %nin% "Other") ▷  
  ggplot(mapping = aes(x = degree,  
                        y = prop,  
                        ymin = prop - 2*prop_se,  
                        ymax = prop + 2*prop_se,  
                        fill = race,  
                        color = race,  
                        group = race))  
  
dodge_w ← position_dodge(width = 0.9)
```

Draw the plot

```
p + geom_col(position = dodge_w, alpha = 0.2) +  
  geom_errorbar(position = dodge_w, width = 0.2) +  
  scale_x_discrete(labels = wrap_format(10)) +  
  scale_y_continuous(labels = label_percent()) +  
  scale_color_brewer(type = "qual",  
                     palette = "Dark2") +  
  scale_fill_brewer(type = "qual",  
                     palette = "Dark2") +  
  labs(title = "Educational Attainment by Race",  
       subtitle = "GSS 1976-2016",  
       fill = "Race",  
       color = "Race",  
       x = NULL, y = "Percent") +  
  facet_wrap(~ year, ncol = 2)
```



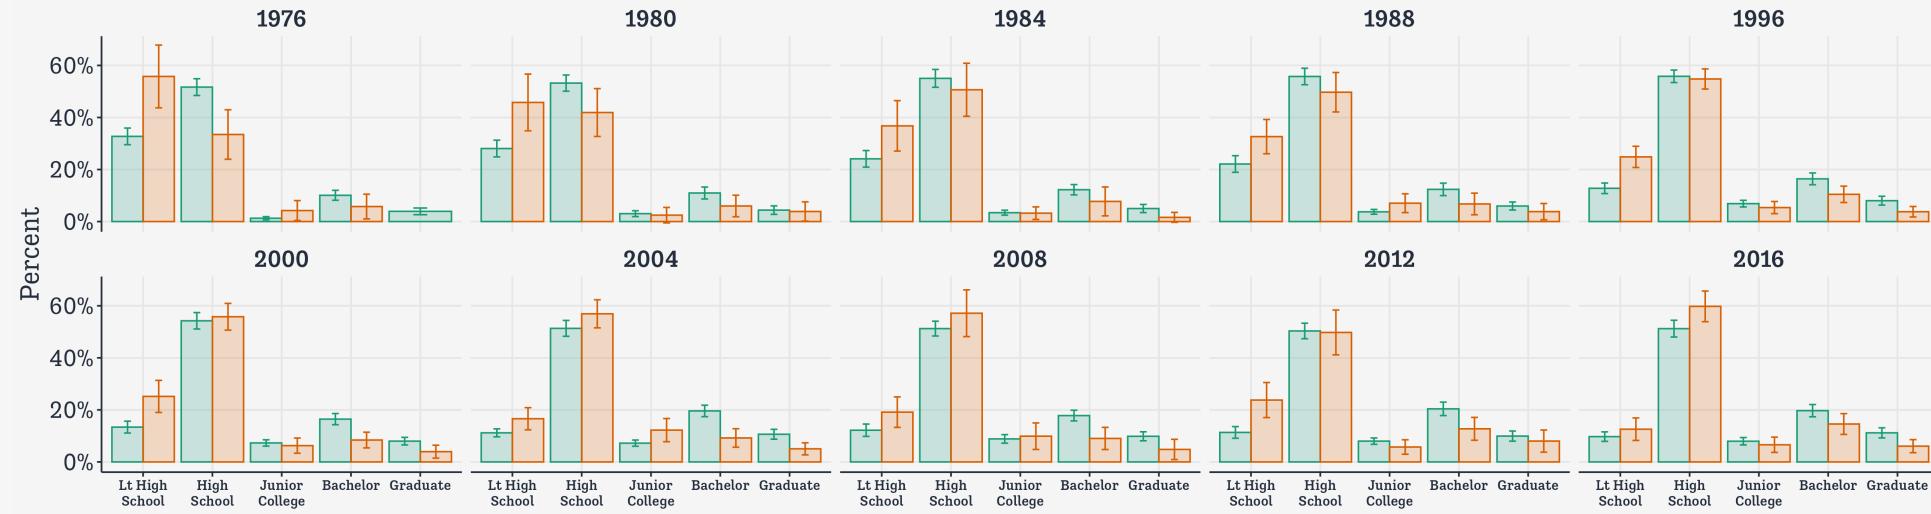
In full (but switch to rows)

```
p_out ← p +
  geom_col(position = dodge_w, alpha = 0.2) +
  geom_errorbar(position = dodge_w, width = 0.2) +
  scale_x_discrete(labels = wrap_format(10)) +
  scale_y_continuous(labels = label_percent()) +
  scale_color_brewer(type = "qual",
                      palette = "Dark2") +
  scale_fill_brewer(type = "qual",
                     palette = "Dark2") +
  labs(title = "Educational Attainment by Race",
       subtitle = "GSS 1976-2016",
       fill = "Race",
       color = "Race",
       x = NULL, y = "Percent") +
  facet_wrap(~ year, nrow = 2) +
  theme(axis.text.x =
        element_text(size = rel(0.6),
                     face = "bold"))
```

Educational Attainment by Race

GSS 1976-2016

Race █ White █ Black



Is this figure
effective? Not
really!

Let's try a different view

```
p ← out_yrd ▷  
  drop_na() ▷  
  filter(race %nin% "Other",  
         degree %nin% "Junior College") ▷  
  ggplot(mapping = aes(x = year, y = prop,  
                        ymin = prop - 2*prop_se,  
                        ymax = prop + 2*prop_se,  
                        fill = race, color = race,  
                        group = race))  
  
p_out ← p +  
  geom_ribbon(mapping = aes(color = NULL),  
              alpha = 0.3) +  
  geom_line(linewidth = rel(1.25)) +  
  scale_y_continuous(labels = label_percent()) +  
  scale_color_brewer(type = "qual", palette = "Dark2") +  
  scale_fill_brewer(type = "qual", palette = "Dark2") +  
  facet_wrap(~ degree, ncol = 2) +  
  labs(title = "Educational Attainment by Race",  
       subtitle = "GSS 1976-2016", fill = "Race",  
       color = "Race", x = NULL, y = "Percent")
```

Educational Attainment by Race

GSS 1976-2016

Race ■ White ■ Black

